

Session 3 - Descriptive Statistics

R Training

María Reyes Retana

The World Bank | January 2025



Government Analytics and R Training:

Strengthening Public Sector Reporting and Data Analysis

January 13 – January 17, 2025



Introduction

Introduction

- We learned yesterday how to conduct statistical programming and export the results in `.csv` files
- However, sometime we might need more refined tables than simple (and ugly) csvs



Introduction

- That's what today's session is about, along with an introduction of the pipe (`%>%`)



Introduction

Relevance to your work

In your **annual reports**, you often include:

- Summary tables like **revenue by year**, **spending across categories**, or **staff counts by department**.
- These tables help summarize key trends and patterns for decision-makers.

Table 8: Programmed PSRL for 2023 (In GH¢)

Month	Diesel	LPG	Petrol	MGO-F	Gas Oil	Gas Oil	Unified/ Naphtha	Total
Jan-23	23,872,267.72	3,682,560.43	31,661,429.70	55,898.08	5,237,078.13	753,524.51	18,461.87	65,281,220.45
Feb-23	24,309,842.16	3,182,459.63	28,976,008.73	50,014.07	4,427,220.69	674,206.14	24,212.29	61,643,963.71
Mar-23	24,096,033.18	3,818,951.55	28,386,586.25	48,248.87	4,211,258.70	650,410.63	24,212.29	61,235,701.47
Apr-23	26,358,863.90	3,409,778.17	31,439,960.64	46,483.67	4,049,287.21	626,615.12	30,265.37	65,961,254.09
May-23	24,599,710.89	3,500,705.59	31,281,630.87	44,130.07	4,103,277.71	594,887.77	24,212.29	64,148,555.19
Jun-23	27,312,455.37	3,909,878.97	31,923,406.89	43,541.66	3,995,296.72	586,955.93	24,212.29	67,795,747.84
Jul-23	27,123,060.55	4,091,733.81	32,816,321.77	42,364.86	3,887,315.73	571,092.26	24,212.29	68,556,101.27
Aug-23	25,959,652.42	4,000,806.39	31,640,924.49	41,188.06	4,265,249.20	555,228.58	24,212.29	66,487,261.44
Sep-23	21,590,679.66	4,000,806.39	33,402,163.23	52,367.68	4,589,192.18	705,933.48	27,238.83	64,368,381.44
Oct-23	27,235,725.40	3,909,878.97	27,364,475.58	52,956.08	4,697,173.17	713,865.32	27,238.83	64,001,313.35
Nov-23	28,242,906.95	3,909,878.97	36,501,774.26	54,132.88	5,129,097.14	729,728.99	26,936.18	74,594,455.37
Dec-23	27,965,119.88	4,046,270.10	39,826,295.00	57,074.88	5,399,049.62	769,388.18	27,238.83	78,090,436.49
Total	308,666,318.08	45,463,708.96	385,220,977.42	588,400.87	53,990,496.19	7,931,836.90	302,653.68	802,164,392.11

Source: NPA

Billed PSRL for 2023

- Today, we will practice **creating similar summary tables** using our staff list data.
- While we're using a simple dataset today, the same steps can be applied to your own data for reports with many more variables.

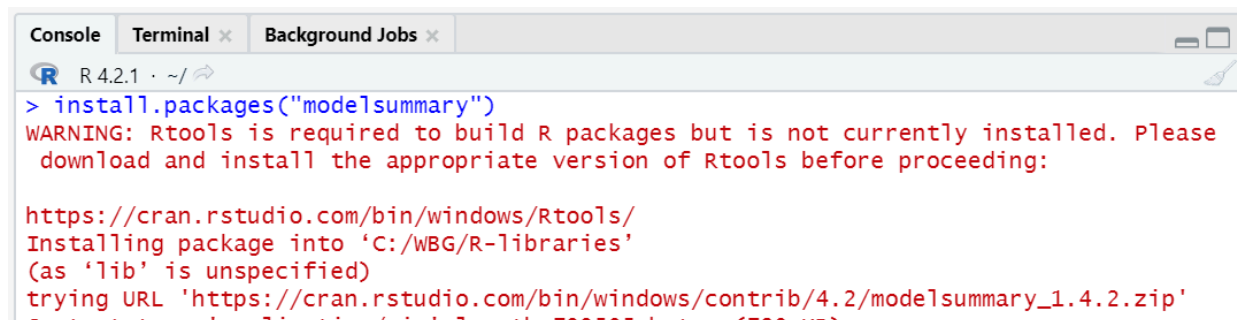
Introduction

Exercise 1a: Getting the packages for today's session

We're going to use several R packages in this session: `modelsummary`, `huxtable` and `dplyr` (you already have this). You can download the solutions 3 from the materials package, where I listed all the packages you need to install, and the instructions. You can start with the script as your base for today's exercise.

1. Install needed packages: (the package `modelsummary` has many dependencies)

```
install.packages("modelsummary")  
install.packages("huxtable")  
install.packages("ggplot2")  
install.packages("rstudioapi")  
install.packages("openxlsx")  
install.packages("dplyr")
```

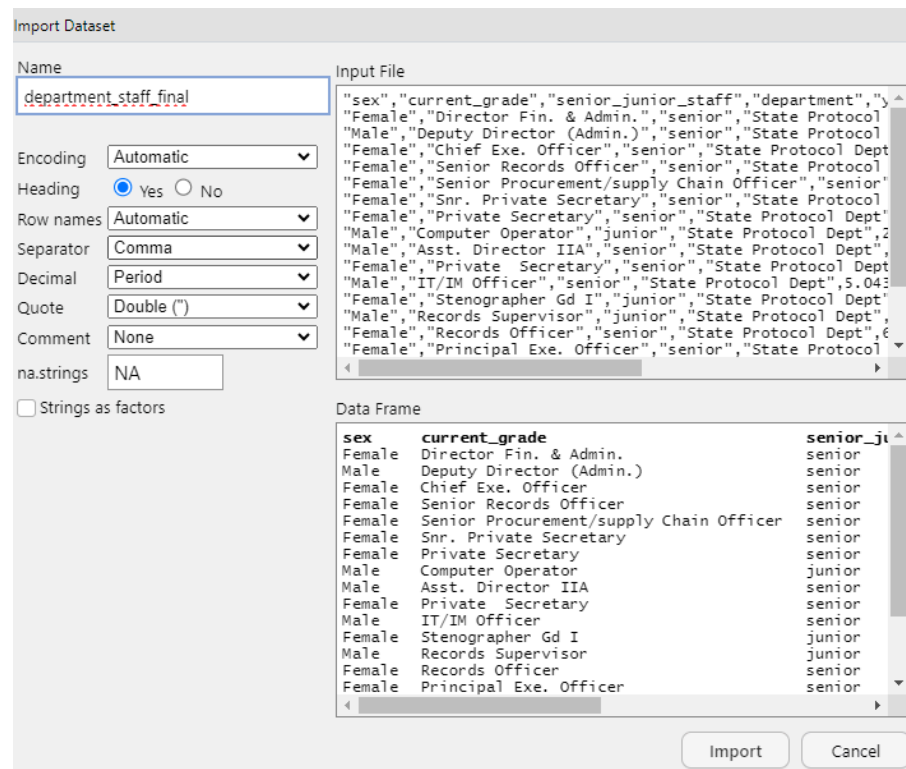


```
Console Terminal x Background Jobs x  
R 4.2.1 ~/  
> install.packages("modelsummary")  
WARNING: Rtools is required to build R packages but is not currently installed. Please  
download and install the appropriate version of Rtools before proceeding:  
  
https://cran.rstudio.com/bin/windows/Rtools/  
Installing package into 'C:/WBG/R-libraries'  
(as 'lib' is unspecified)  
trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.2/modelsummary_1.4.2.zip'  
Warning: package 'modelsummary' is not available for this version of R  
Warning: installing from source is not supported on this platform
```

Introduction

Exercise 1b: Download and load the data we'll use

1. Go to <https://osf.io/rk9ht> and download the file
2. In RStudio, go to **File** > **Import Dataset** > **From Text (base)** and select the file `department_staff_final.csv`
 - If you don't know where the file is, remember to check in your **Downloads** folder
3. Select **Import**



Introduction

You should have one dataframe loaded in the environment after this.

Environment	History	Connections	Git	Tutorial
<div><div> </div><div> Import Dataset ▾</div><div> 195 MiB ▾</div><div></div></div>				
<div><div>R ▾</div><div> Global Environment ▾</div></div>				
Data				
<div><div> department_staff_final</div></div>			8645 obs. of 6 variables	

Introduction

Recap: always know your data!

- This is the data that we created yesterday! (after merging the age data)

```
glimpse(department_staff_final)
```

```
## Rows: 8,645
```

```
## Columns: 6
```

```
## $ sex                <chr> "Female", "Male", "Female", "Female", "Female", "F...
```

```
## $ current_grade      <chr> "Director Fin. & Admin.", "Deputy Director (Admin....
```

```
## $ senior_junior_staff <chr> "senior", "senior", "senior", "senior", "senior", ...
```

```
## $ department         <chr> "State Protocol Dept", "State Protocol Dept", "Sta...
```

```
## $ years_of_service    <dbl> 30.220397, 16.049281, 33.968515, 11.553730, 21.180...
```

```
## $ age                 <dbl> 60.15058, 41.74949, 59.43053, 50.55989, 51.28268, ...
```

Piping

The pipe (%>%) operator

- Before diving into the contents of today's session, we will to introduce a **super useful tool**: **the pipe (%>%)**.
- The pipe is part of the `dplyr` package. It helps to write code in a way that is easier to read and understand.
- Reading and understanding multiple operations can be difficult.
- The pipe operator (%>%) can help with this.

The pipe operator

- With the pipe, code reads from left to right, top to bottom, which is more intuitive.
- `%>%` can be read as "then" and simplifies code structure.

For example let's see at this mock get to work sequence example:

Without pipe (`%>%`) (hard to read and understand the sequence)

```
go_to_work(make_breakfast(work_out(brush_teeth(wake_up(Mer)))))
```

With pipe (`%>%`) (the order is clear)

```
Mer %>%  
  wake_up() %>%  
  brush_teeth() %>%  
  work_out() %>%  
  make_breakfast() %>%  
  go_to_work()
```

The pipe

As we saw yesterday the functions are normally organized around a set of verbs, or actions to be taken.

The pipe

As we saw yesterday the functions are normally organized around a set of verbs, or actions to be taken.

- Most *verbs* work as follows:

`verb(data.frame, what to do)`
1st argument 2nd argument

The pipe

As we saw yesterday the functions are normally organized around a set of verbs, or actions to be taken.

- Most *verbs* work as follows:

$\text{verb}(\underbrace{\text{data.frame}}_{\text{1st argument}}, \underbrace{\text{what to do}}_{\text{2nd argument}})$

- Alternatively you can (**should**) use the **pipe** operator `%>%`:

$\underbrace{\text{data.frame}}_{\text{1st argument}} \underbrace{\text{\%>\%}}_{\text{"pipe" operator}} \text{verb}(\underbrace{\text{what to do}}_{\text{2nd argument}})$

We will start using the pipe from this point. Please ask if something is not clear.

Tip 💡 : Use Shift + Ctrl/Cmd + M as a shortcut for the pipe operator.

The pipe

- You probably remember this piece of code from one of yesterday's exercise:

```
# Filter only female employees:  
temp1 <- filter(department_staff_list, sex == "Female") # filter by female  
  
# Sort previous result by years of service  
department_female <- arrange(temp1, years_of_service) # order by years of service
```

The pipe

This code works, but the problem with it is that it makes us generate unnecessary intermediate dataframes (`temp1`) that store results temporarily

R ▾ Global Environment ▾	
Data	
▶ department_female	3559 obs. of 6 variables
▶ department_staff_age	8714 obs. of 2 variables
▶ department_staff_list	8714 obs. of 6 variables
▶ temp1	3559 obs. of 6 variables

The pipe

Instead, we can use pipes to **pass the results of a function and apply a new function on top of it** (just like Mer's waking up sequence)

```
# Filter only female employees:  
temp1 <- filter(department_staff_list, sex == "Female")  
  
# Sort previous result by years of service  
department_female <- arrange(temp1, years_of_service)
```

```
# The same but with pipes:  
department_female <- filter(department_staff_list, sex == "Female") %>%  
  arrange(years_of_service)
```

- The usefulness of the pipe (%>%) becomes more evident when the code starts to get more complicated.

Piping

```
# Filter only female employees:  
temp1 <- filter(department_staff_list, sex == "Female")  
  
# Sort previous result by years of service  
department_female <- arrange(temp1, years_of_service)
```

```
# The same but with pipes:  
department_female <- filter(department_staff_list, sex == "Female") %>%  
  arrange(years_of_service)
```

There are several important details to notice here:

- 1.- The resulting dataframe `department_female` is **the same in both cases**

Piping

```
# Filter only female employees:  
temp1 <- filter(department_staff_list, sex == "Female")  
  
# Sort previous result by years of service  
department_female <- arrange(temp1, years_of_service)
```

```
# The same but with pipes:  
department_female <- filter(department_staff_list, sex == "Female") %>%  
  arrange(years_of_service)
```

2.- Notice that the functions `arrange()` and `filter()` used after the pipes now have only **one argument instead of two**. This is because when using pipes the first argument is implied to be result of the function before the pipes

Exercise 2: filtering and sorting revisited

1. Apply the same filtering and sorting now with pipes

Solution

```
department_female <- filter(department_staff_final, sex == "Female") %>%  
  arrange(years_of_service)
```

- Note: that our dataframe now is named `depatment_staf_final` (is the joined dataframe from yesterday's session)

Piping

Now we will not have any annoying intermediate results stored in our environment!

- Good code is code that is both correct (does what it's supposed to) and it's easy to understand
- Piping is **instrumental for writing good code in R**

Always use pipes!

Now that you now about the power of the pipes, use them wisely!

- Remember that pipes are part of the library `dplyr`, you need to load it before using them
- Pipes also improve code clarity drastically
- Many R coders use pipes and internet examples assume you know them
- **We'll use pipes now in the next examples and exercises of the rest of this training**



Quick summary statistics

Quick Summary Statistics

Refresher: Grouping and Summarizing Data

Yesterday, we learned how to:

1. **Group data** using `group_by()`
2. **Summarize** results with `summarise()`

This is a powerful tool to create **summary tables**—such as totals, averages, or counts—that are essential for your annual reports.

Example

```
# Summarize total revenue by month
summary_table <- psrl_data %>%
  group_by(Month) %>%
  summarise(Total_Revenue = sum(Revenue, na.rm = TRUE))

print(summary_table)
```

Quick summary statistics

Applying This to Your Work

Annual Report Table

Table 8: Programmed PSRL for 2023 (In GH¢)

Month	Diesel	LPG	Petrol	MGO-F	Gas Oil	Gas Oil	Unified/ Naphtha	Total
Jan-23	23,872,267.72	3,682,560.43	31,661,429.70	55,898.08	5,237,078.13	753,524.51	18,461.87	65,281,220.45
Feb-23	24,309,842.16	3,182,459.63	28,976,008.73	50,014.07	4,427,220.69	674,206.14	24,212.29	61,843,963.71
Mar-23	24,096,033.18	3,818,951.55	28,386,586.25	48,248.87	4,211,258.70	650,410.63	24,212.29	61,235,701.47
Apr-23	26,358,863.90	3,409,778.17	31,439,960.64	46,483.67	4,049,287.21	626,615.12	30,265.37	65,961,254.09
May-23	24,599,710.89	3,500,705.59	31,281,630.87	44,130.07	4,103,277.71	594,887.77	24,212.29	64,148,555.19
Jun-23	27,312,455.37	3,909,878.97	31,923,406.89	43,541.66	3,995,296.72	586,955.93	24,212.29	67,795,747.84
Jul-23	27,123,060.55	4,091,733.81	32,816,321.77	42,364.86	3,887,315.73	571,092.26	24,212.29	68,556,101.27
Aug-23	25,959,652.42	4,000,806.39	31,640,924.49	41,188.06	4,265,249.20	555,228.58	24,212.29	66,487,261.44
Sep-23	21,590,679.66	4,000,806.39	33,402,163.23	52,367.68	4,589,192.18	705,933.48	27,238.83	64,368,381.44
Oct-23	27,235,725.40	3,909,878.97	27,364,475.58	52,956.08	4,697,173.17	713,865.32	27,238.83	64,001,313.35
Nov-23	28,242,906.95	3,909,878.97	36,501,774.26	54,132.88	5,129,097.14	729,728.99	26,936.18	74,594,455.37
Dec-23	27,965,119.88	4,046,270.10	39,826,295.00	57,074.88	5,399,049.62	769,388.18	27,238.83	78,090,436.49
Total	308,666,318.08	45,463,708.96	385,220,977.42	588,400.87	53,990,496.19	7,931,836.90	302,653.68	802,164,392.11

Source: NPA

PSRL for 2023

Monthly totals for different products.

Code to Recreate the Table

```
# Summarize all numeric columns by Month
summary_table <- psrl_data %>%
  group_by(Month) %>%
  summarise(across(where(is.numeric), \(x) sum(x, na.rm = TRUE)))
```

Quick summary statistics

```
# Print the summary table  
print(summary_table)
```

```
## # A tibble: 12 × 8  
##   Month Diesel LPG Petrol MGO_F Gas_Oil Unified_Naphtha Total  
##   <chr>   <dbl> <dbl>  <dbl> <dbl>   <dbl>      <dbl> <dbl>  
## 1 Apr-23   140    70   115    40    35         22   422  
## 2 Aug-23   170    85   135    48    45         30   513  
## 3 Dec-23   200   100   155    60    55         40   610  
## 4 Feb-23   130    65   105    35    30         18   383  
## 5 Jan-23   120    60   100    30    25         15   350  
## 6 Jul-23   160    80   130    50    42         28   490  
## 7 Jun-23   155    78   125    45    40         26   469  
## 8 Mar-23   125    63   110    33    28         20   379  
## 9 May-23   150    75   120    42    38         24   449  
## 10 Nov-23  190    95   150    58    52         38   583  
## 11 Oct-23  180    90   145    55    50         35   555  
## 12 Sep-23  165    83   140    47    48         32   515
```

Quick summary statistics

Beyond Basic Summaries: Customizing Your Results

We learned yesterday how to produce dataframes with results and export them.

But what if you want to ... ?

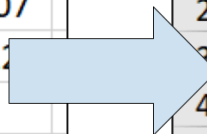
- ...export results in a different format (example: Excel)
- ...further customize which rows and columns to display in a result
- ...format the results you export

Quick summary statistics

You will need `modelsummary` and `huxtable` for this

- These libraries allow you to export results in a customized way
- We chose a combination of both because together they export a large range of output types and allow fine-grained customization of outputs

	A	B	C	D	E	F
1		mean	sd	min	median	max
2	modified_id	#####	#####	19832	5008712	1.2E+07
3	taxperiod	201907	3	201901	201907	201911
4	age	14	8	1	13	30
5	income	3284	8242	0	907	139395
6	vat_liability	591	1484	0	163	25091



	A	B	C	D	E	
1		Unique (#)	Missing (%)	Mean	SD	
2	modified_id	984	0	5448915	3758602	1
3	taxperiod	12	0	201907	3	2
4	age	30	0	14	8	1
5	income	721	0	3284	8242	0
6	vat_liability	721	0	591	1484	0

Quick summary statistics

We'll start by introducing the function `datasummary_skim()` from `modelsummary`

```
datasummary_skim(data, output, ...)
```

- **data:** the data set to be summarized, the only required argument
- **output:** the type of output desired
- **...:** additional options allow for formatting customization, such as including notes and titles

For example:

```
datasummary_skim(  
  data,  
  output = "default",  
  type = "numeric",  
  title = NULL,  
  notes = NULL,  
  ...  
)
```

Quick summary statistics



Exercise 3: Calculate quick summary statistics

1. Load `modelsummary` with `library(modelsummary)`
2. Use `datasummary_skim()` to create a descriptive statistics table for `department_staff`

```
datasummary_skim(department_staff)
```


Quick summary statistics

You should be seeing this result in the lower right panel of RStudio.

	Unique	Missing Pct.	Mean	SD	Min	Median	Max	Histogram
years_of_service	2916	0	14.7	9.6	0.2	13.7	50.7	
age	5998	0	44.2	8.8	17.5	43.5	66.7	
		N						
sex	Female	3529	40.8					
	Male	5116	59.2					
senior_junior_staff	junior	2408	27.9					
	senior	6237	72.1					
department	Births & Deaths	165	1.9					
	Bureau of Languages	58	0.7					
	Co-operatives	268	3.1					
	Controller & Accountant General	3776	43.7					
	Department of Comm Devt	109	1.3					
	Dept of Children	99	1.1					
	Dept of Gender	100	1.2					
	Dept of Labour	290	3.4					
	Factories Inspectorate	121	1.4					

Quick summary statistics

- Most functions of `modelsummary` summarize only numeric variables by default
- To summarize categorical variables, use the argument `type = "categorical"`

```
datasummary_skim(department_staff, type = "categorical")
```

Quick summary statistics

		N	%
sex	Female	3529	40.8
	Male	5116	59.2
senior_junior_staff	junior	2408	27.9
	senior	6237	72.1
department	Births & Deaths	165	1.9
	Bureau of Languages	58	0.7
	Co-operatives	268	3.1
	Controller & Accountant General	3776	43.7
	Department of Comm Devt	109	1.3
	Dept of Children	99	1.1
	Dept of Gender	100	1.2
	Dept of Labour	290	3.4
	Factories Inspectorate	121	1.4
	Food & Drug	102	1.2

Quick summary statistics

- `datasummary_skim()` is convenient because it's fast, easy, and shows a lot of information
- But what if we wanted to customize what to show? that's when we use `datasummary()` instead, also from the library `modelsummary`

Customized summary statistics

Customized summary statistics

`datasummary()` is very similar to `data_summary_skim()`. The only difference is that it requires a **formula argument**.

`datasummary(formula, data, output)`

- **formula:** a two-sided formula to describe the table as: rows ~ columns
- **data:** the data set to be summarized
- **output:** the type of output desired
- **...**: additional options allow for formatting customization

```
datasummary(  
  var1 + var2 + var3 ~ stat1 + stat2 + stat3 + stat4,  
  data = data  
)
```

Customized summary statistics

Exercise 4:

Create a summary statistics table showing the number of observations, mean, minimum, and maximum for variables `years_of_service` of the dataframe `department_staff`

1. Use `datasummary()` for this:

```
datasummary(  
  years_of_service ~ N + Mean + Min + Max,  
  department_staff  
)
```

	N	Mean	SD	Min	Max
years_of_service	8645	14.73	9.59	0.16	50.66

Customized summary statistics

```
datasummary(  
  years_of_service ~ N + Mean + SD + Min + Max, # this is the formula  
  department_staff ~ # this is the data  
)
```

Some notes:

- The arguments **formula** and **data** are mandatory for `datasummary()`
- All other arguments are optional (like `title = *some-title*`, to add a table title)
- The formula should always be defined as: rows ~ columns
- The rows and columns in the formula are separated by a plus (+) sign

In Excel  you would need to calculate each of the statistics in a new table, by selecting the data and using the appropriate formula.

Customized summary statistics

```
datasummary(  
  years_of_service ~ N + Mean + SD + Min + Max, # this is the formula  
  department_staff      # this is the data  
)
```

In this exercise we used the statistics N (number of observations), mean, SD (standard deviation), Min (minimum), and Max (maximum). Other statistics you can include are:

Statistic	Keyword
Median	Median
25th percentile	P25
75th percentile	P75
In general: percentile XX	PXX
Small histogram	Histogram

Exporting tables

Exporting tables

Remember that both `datasummary_skim()` and `datasummary()` have an optional argument named *output*? We can use it to specify a file path for an output file.

For example:

```
datasummary_skim(department_staff,  
                  output = "quick_stats.docx")
```

Will export the result to the `Documents` folder (in Windows) in a Word file named `quick_stats.docx`

Note for this code to work we would need to install an extra package `pandoc`

Exporting tables

The file type of the output is dictated by the file extension. For example:

File name	File extension	Output type
"quick_stats.docx"	<code>.docx</code>	Word
"quick_stats.pptx"	<code>.pptx</code>	Power Point
"quick_stats.html"	<code>.html</code>	HTML (to open in a web explorer)
"quick_stats.tex"	<code>.tex</code>	Latex
"quick_stats.md"	<code>.md</code>	Markdown

Noticed that we're missing Excel?

Exporting tables

That's because the functions of `modelsummary` can't export to Excel

- Nonetheless, we can use the library `huxtable` as an intermediary to transform results from `modelsummary` functions to Excel files
- `huxtable` is a package for exporting tables in general that allows you to **customize the output you're exporting**
- We'll know how to use it in the next exercise

Exporting tables













Exercise 5: Export a table to Excel

1. Load `huxtable` with `library(huxtable)` (we already did this at the beginning of the session)
2. Run the following code to export the result of `datasummary_skim()` to Excel:

```
# Store the table in a new object  
stats_table <- datasummary_skim(department_staff, output = "huxtable")  
  
# Export this new object to Excel with quick_xlsx()  
quick_xlsx(stats_table, file = "quick_stats.xlsx")
```

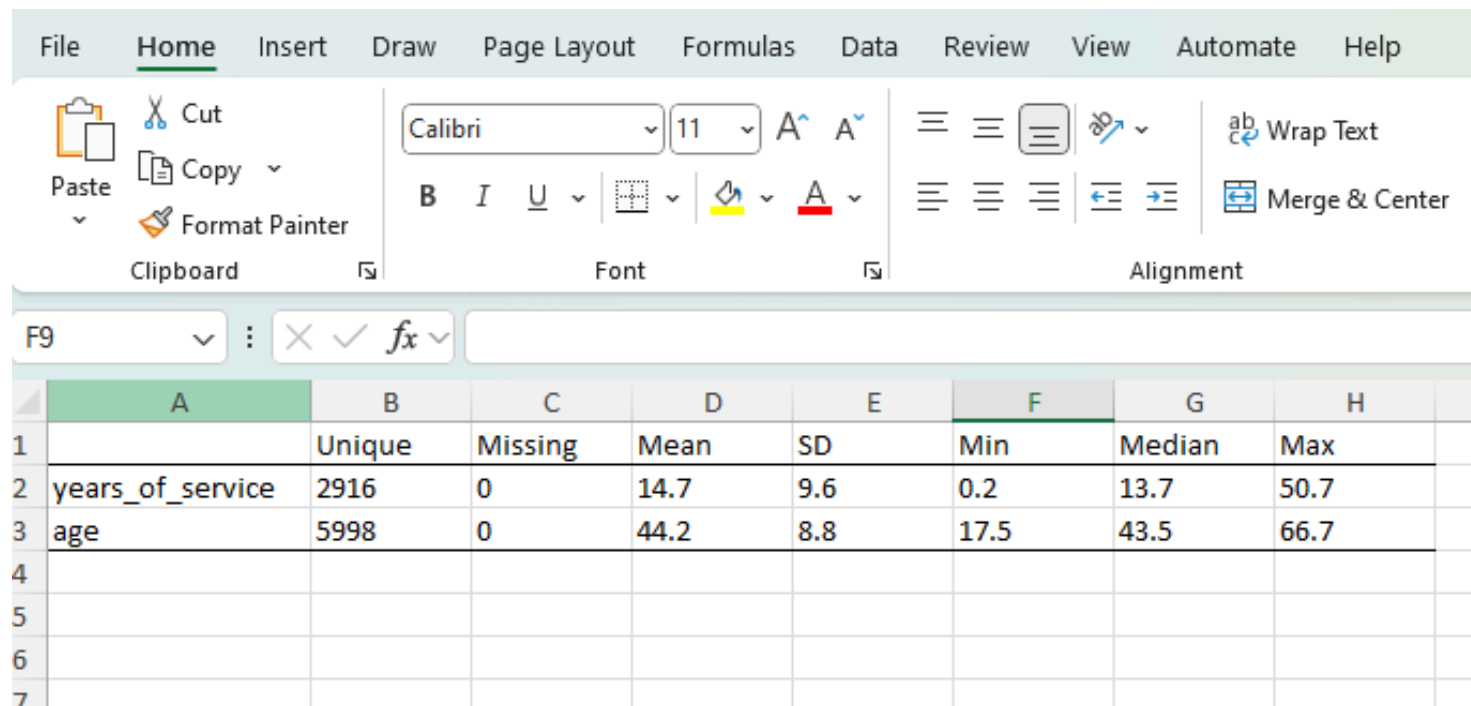
Exporting tables

Now the result will show in your **Documents** folder

 3-descriptive-statistics.html	14/12/2024 3:50 pm
 3-descriptive-statistics.pdf	06/11/2024 4:53 pm
 3-descriptive-statistics.Rmd	14/12/2024 3:50 pm
 4-data-visualization.html	06/11/2024 4:53 pm
 4-data-visualization.pdf	06/11/2024 4:53 pm
 4-data-visualization.Rmd	13/12/2024 5:39 pm
 202410.Rproj	14/12/2024 3:26 pm
 exercises-session1.R	17/11/2024 3:18 pm
 exercises-session2.R	19/11/2024 3:25 pm
 exercises-session3.R	14/12/2024 3:24 pm
 exercises-session4.R	06/11/2024 4:53 pm
 quick_stats.xlsx	14/12/2024 3:21 pm

Exporting tables

And you can open it with Excel for further customization if you want...



The screenshot shows the Microsoft Excel interface with the 'Home' ribbon selected. The ribbon includes tabs for File, Home, Insert, Draw, Page Layout, Formulas, Data, Review, View, Automate, and Help. The 'Home' ribbon is divided into three main groups: Clipboard (with Cut, Copy, Paste, and Format Painter), Font (with font face, size, bold, italic, underline, color, and background color), and Alignment (with text alignment, orientation, and Merge & Center). Below the ribbon, the formula bar shows 'F9' and a function icon. The spreadsheet grid displays a table with 8 columns (A-H) and 7 rows. The table contains summary statistics for 'years_of_service' and 'age'.

	A	B	C	D	E	F	G	H
1		Unique	Missing	Mean	SD	Min	Median	Max
2	years_of_service	2916	0	14.7	9.6	0.2	13.7	50.7
3	age	5998	0	44.2	8.8	17.5	43.5	66.7
4								
5								
6								
7								

- However... remember that any manual changes will be hard to track affecting the reproducibility of your work.

Exporting tables

```
# Store the table in a new object
stats_table <- datasummary_skim(department_staff, output = "huxtable")

# Export this new object to Excel with quick_xlsx()
quick_xlsx(stats_table, file = "quick_stats.xlsx")
```

Some comments about this code:

- `quick_xlsx()` is a function from `huxtable`. The first argument is the object we export and the second is the file name. We could also use a file path here
- Note that we now use the argument `output = "huxtable"` in `datasummary_skim()`. This tells R that the output should be an object type that we can operate later with `huxtable` functions, such as `quick_xlsx()`

Customizing table outputs

Customizing table outputs

The code below shows how the table `stats_table` can be formatted:

```
# We start with stats_table:
stats_table %>%
  # Use first row as table header
  set_header_rows(1, TRUE) %>%
  # Use first column as row header
  set_header_cols(1, TRUE) %>%
  # Don't round large numbers
  set_number_format(everywhere, 2:ncol(.), "%9.0f") %>%
  # Center cells in first row
  set_align(1, everywhere, "center") %>%
  # Set a theme for quick formatting
  theme_blue()
```

	Unique	Missing Pct.	Mean	SD	Min	Median	Max
years_of_service	2916	0	15	10	0	14	51
age	5998	0	44	9	18	44	67

Customizing table outputs

Exercise 6: Export a customized table to Excel

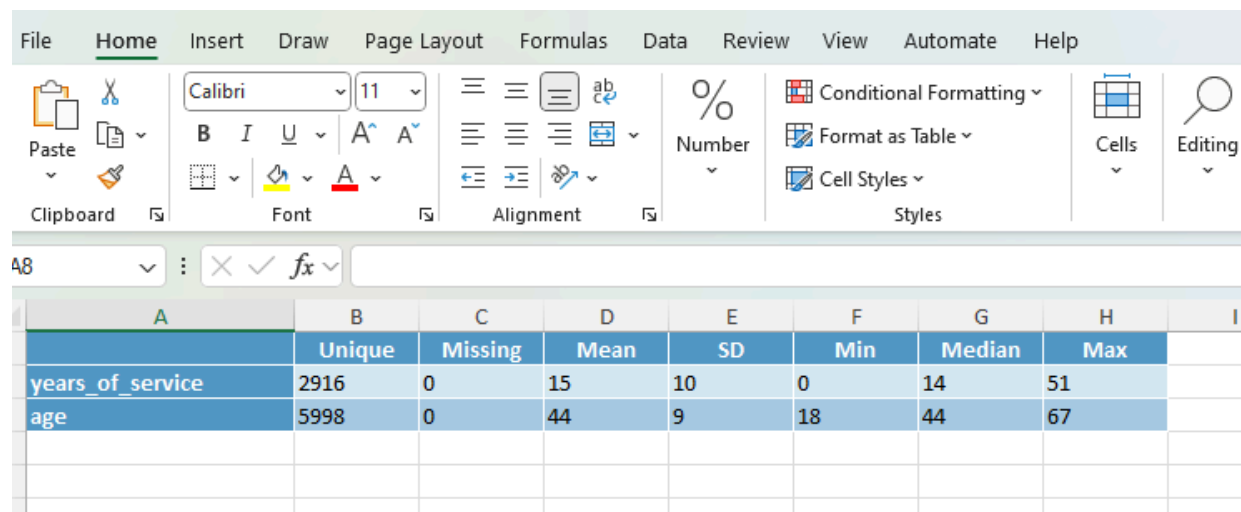
1.- Customize `stats_table` in a new object called `stats_table_custom`

```
stats_table_custom <- stats_table %>%  
  # Use first row as table header  
  set_header_rows(1, TRUE) %>%  
  # Use first column as row header  
  set_header_cols(1, TRUE) %>%  
  # Don't round large numbers  
  set_number_format(everywhere, 2:ncol(.), "%9.0f") %>%  
  # Center cells in first row  
  set_align(1, everywhere, "center") %>%  
  # Set a theme for quick formatting  
  theme_blue()
```

2.- Export `stats_table_custom` to a file named `stats-custom.xlsx` with `quick_xlsx()`

```
quick_xlsx(  
  stats_table_custom,  
  file = "stats-custom.xlsx"  
)
```

Customizing table outputs



The screenshot shows the Microsoft Excel interface with the 'Home' ribbon selected. The ribbon includes groups for Clipboard, Font, Alignment, Number, Styles, Cells, and Editing. Below the ribbon, the formula bar shows 'A8'. The worksheet contains a table with the following data:

	A	B	C	D	E	F	G	H	I
		Unique	Missing	Mean	SD	Min	Median	Max	
years_of_service		2916	0	15	10	0	14	51	
age		5998	0	44	9	18	44	67	

Customizing table outputs

Notice that here in the first part of the exercise we stored the result in a new object

```
stats_table_custom <- stats_table %>% # <---- here
  set_header_rows(1, TRUE) %>%
  set_header_cols(1, TRUE) %>%
  set_number_format(everywhere, 2:ncol(.), "%9.0f") %>%
  set_align(1, everywhere, "center") %>%
  theme_blue()
```

This is the object that we export later with `quick_xlsx()`

```
quick_xlsx(
  stats_table_custom,
  file = "stats-custom.xlsx"
)
```

Customizing table outputs

Before:

File Home Insert Draw Page Layout Formulas Data Review View Automate Help

Clipboard: Paste, Copy, Format Painter

Font: Calibri, 11, Bold, Italic, Underline, Color, Background Color, Text Color

Alignment: Wrap Text, Merge & Center

	A	B	C	D	E	F	G	H
1		Unique	Missing	Mean	SD	Min	Median	Max
2	years_of_service	2916	0	14.7	9.6	0.2	13.7	50.7
3	age	5998	0	44.2	8.8	17.5	43.5	66.7
4								
5								
6								
7								

After:

File Home Insert Draw Page Layout Formulas Data Review View Automate Help

Clipboard: Paste, Copy, Format Painter

Font: Calibri, 11, Bold, Italic, Underline, Color, Background Color, Text Color

Alignment: Wrap Text, Merge & Center

Styles: Conditional Formatting, Format as Table, Cell Styles

	A	B	C	D	E	F	G	H	I
		Unique	Missing	Mean	SD	Min	Median	Max	
	years_of_service	2916	0	15	10	0	14	51	
	age	5998	0	44	9	18	44	67	

Customizing table outputs

We used `theme_blue()`. Other available themes are:

theme_plain

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_basic

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_compact

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_article

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_bright

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_grey

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_blue

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_green

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_mondrian

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_orange

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

theme_stripped

Type	Price	Sugar content
Strawberry	1.90	40.00%
Raspberry	2.10	35.00%
Plum	1.80	50.00%

Use it on your work

Key Takeaways:

- This was a **basic example** with a few variables from your staff list, but the **possibilities are endless**.
- With this and the contents from yesterday's session, you can create **summaries** of **anything you can think of**.

Real-World Example:

Annual Report: Programmed vs. Billings for 2023 (In GH¢)

Table 9: Programmed versus Billings for 2023 (In GH¢)

Month	Programme	Actual Billings	Variance	%Variance
Jan	65,281,220.45	56,009,815.06	(9,271,405.39)	-14.20%
Feb	61,643,963.71	53,401,114.00	(8,242,849.71)	-13.37%
Mar	61,235,701.47	65,176,751.96	3,941,050.49	6.44%
Apr	65,961,254.09	60,430,666.88	(5,530,587.21)	-8.38%
May	64,148,555.19	66,670,601.74	2,522,046.55	3.93%
Jun	67,795,747.84	60,291,190.36	(7,504,557.48)	-11.07%
Jul	68,556,101.27	63,948,059.60	(4,608,041.67)	-6.72%
Aug	66,487,261.44	64,451,574.60	(2,035,686.84)	-3.06%
Sep	64,368,381.44	59,259,571.44	(5,108,810.00)	-7.94%
Oct	64,001,313.35	59,850,932.04	(4,150,381.31)	-6.48%
Nov	74,594,455.37	68,070,927.90	(6,523,527.47)	-8.75%
Dec	78,090,436.49	66,271,765.20	(11,818,671.29)	-15.13%
Total	802,164,392.11	743,832,970.78	(58,331,421.33)	-7.27%

Source: NPA

- If we have the data, you can easily create summaries like this directly in R.
- Once written, the code can be re-used for the next year or quarter.

Questions?

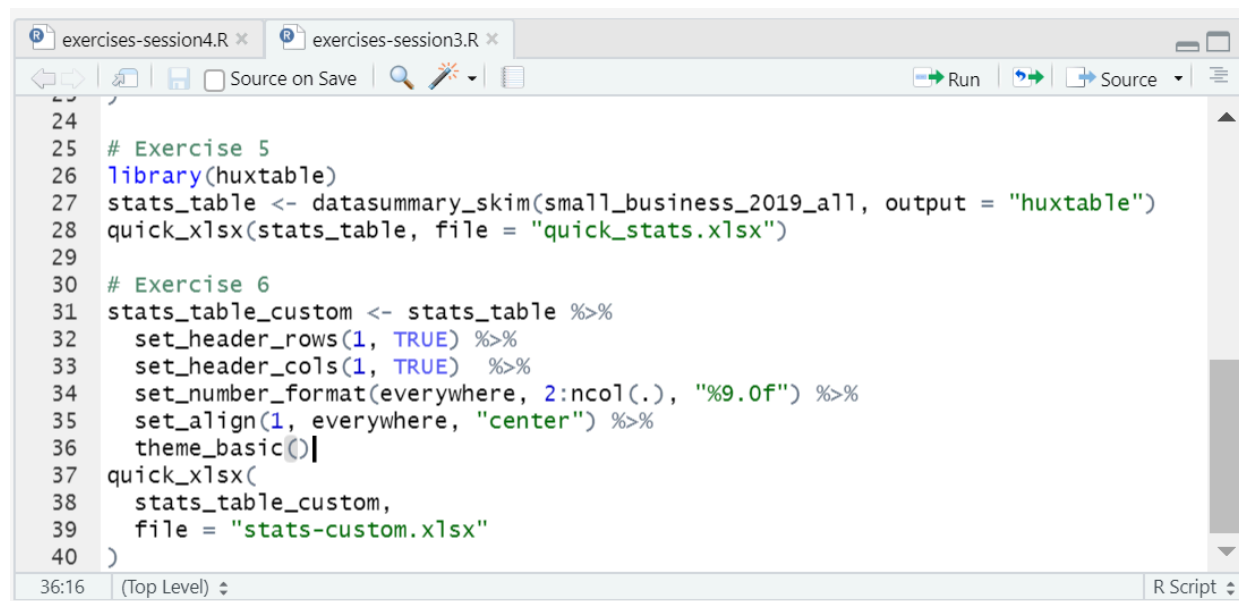


Wrapping up

Wrapping up

Save your work!

Click the floppy disk to save the script you wrote in this session.



The screenshot shows the RStudio interface with two tabs open: 'exercises-session4.R' and 'exercises-session3.R'. The active tab is 'exercises-session4.R', which contains R code for two exercises. The code uses the 'huxtable' package to create and format tables. Exercise 5 creates a table from 'small_business_2019_all' and saves it as 'quick_stats.xlsx'. Exercise 6 takes the 'stats_table' and formats it with custom headers, column widths, and alignment, saving it as 'stats-custom.xlsx'. The status bar at the bottom indicates the cursor is at line 36, column 16, and the file is an 'R Script'.

```
24  
25 # Exercise 5  
26 library(huxtable)  
27 stats_table <- datasummary_skim(small_business_2019_all, output = "huxtable")  
28 quick_xlsx(stats_table, file = "quick_stats.xlsx")  
29  
30 # Exercise 6  
31 stats_table_custom <- stats_table %>%  
32   set_header_rows(1, TRUE) %>%  
33   set_header_cols(1, TRUE) %>%  
34   set_number_format(everywhere, 2:ncol(.), "%9.0f") %>%  
35   set_align(1, everywhere, "center") %>%  
36   theme_basic() |  
37   quick_xlsx(  
38     stats_table_custom,  
39     file = "stats-custom.xlsx"  
40   )
```

Wrapping up

What else is available?

- This was a short overview of how `modelsummary` and `huxtable` work together to produce professional-looking table outputs in R
- Other formatting options are: (all from `huxtable`)

Formatting	Command
Export in new Excel tabs instead of new files	<code>as_Workbook()</code>
Change row names	<code>add_rownames()</code>
Change column names	<code>add_colnames()</code>
Cells in bold	<code>set_bold()</code>
Cells in italics	<code>set_italic()</code>
Cell font size	<code>font_size()</code>
Cell color	<code>background_color()</code>

Wrapping up

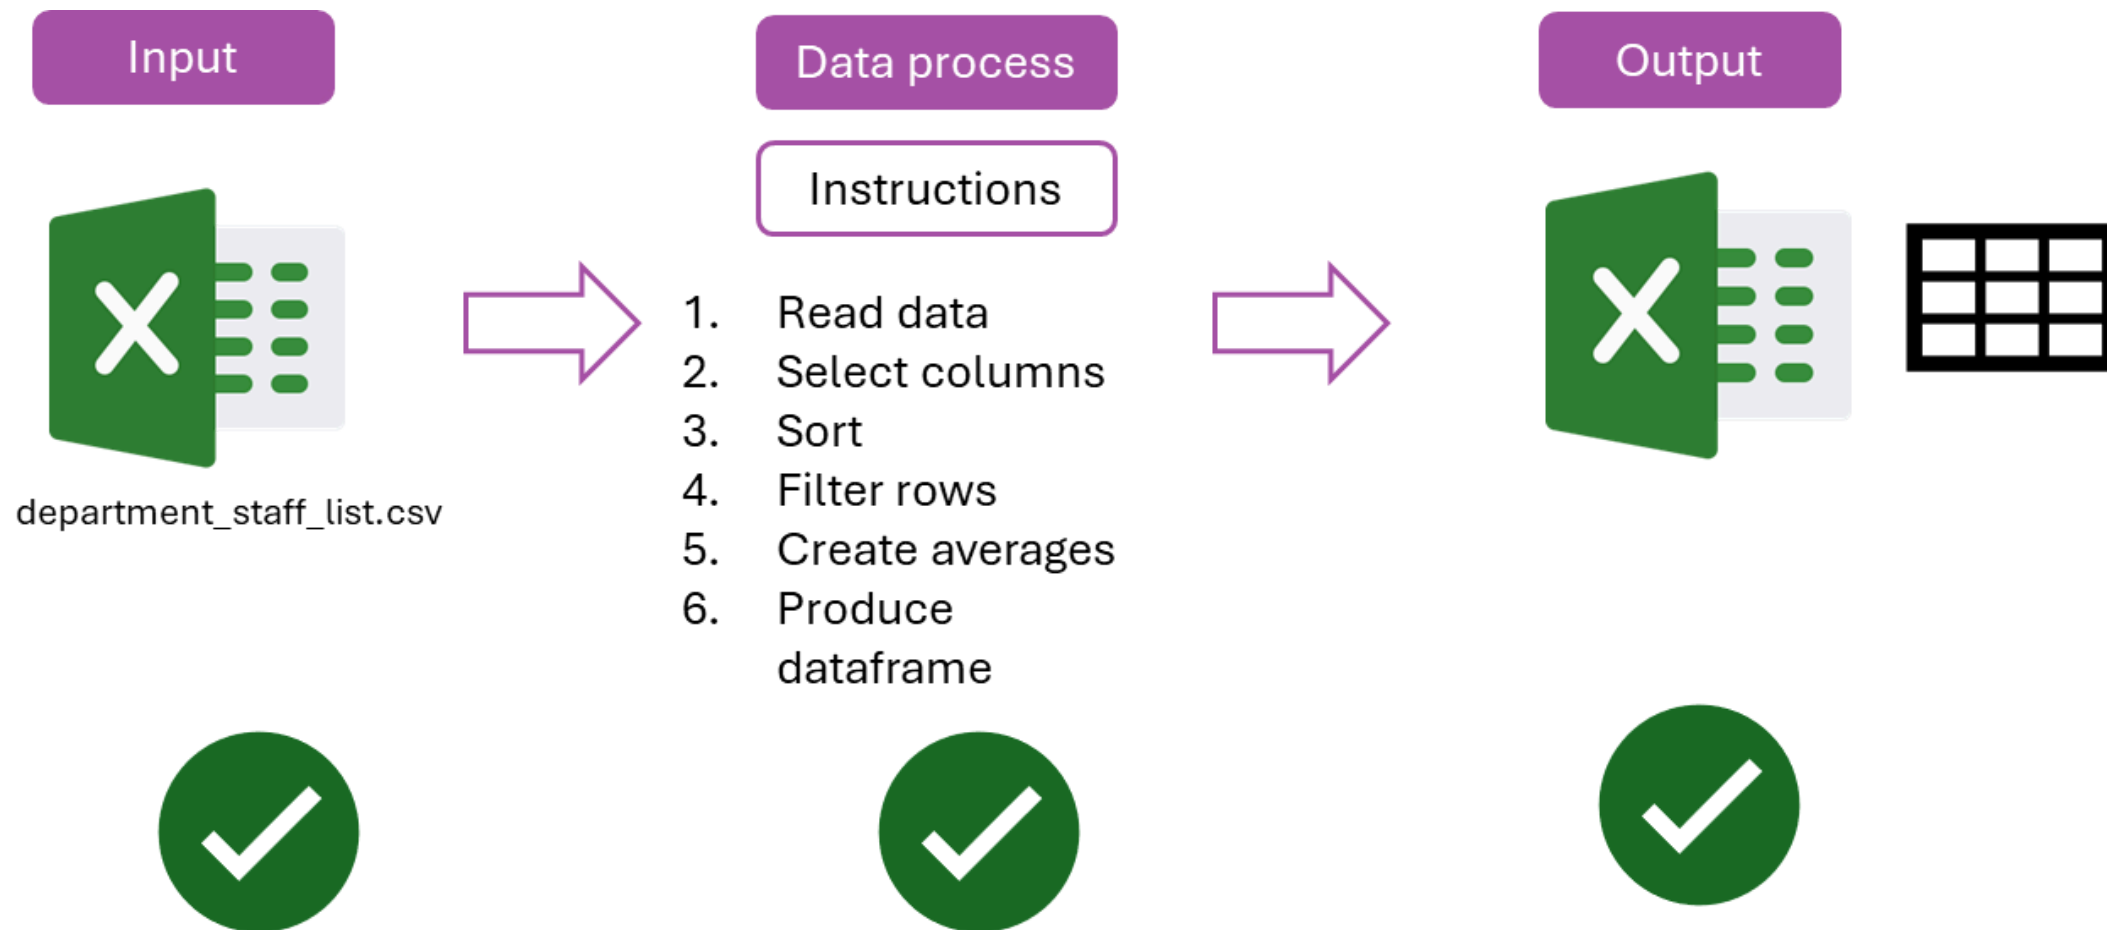
What else is available?

More of this is explained in the libraries documentation:

- `modelsummary` documentation: <https://modelsummary.com/index.html>
- `huxtable` documentation: <https://hughjonesd.github.io/huxtable/>

Wrapping up

This session



Wrapping up

Next session (last one)

Input




department_staff_list.csv

Data process

Instructions

1. Read data
2. Select columns
3. Sort
4. Filter rows
5. Create averages
6. Produce dataframe

Output



Thanks! // ¡Gracias! // Obrigado!

