

# Session 6 - Spatial Data

R for Stata Users

---

Luiza Andrade, Rob Marty, Rony Rodriguez-Ramirez, Luis Eduardo San Martin, Leonardo Viotti

The World Bank – DIME | [WB Github](#)

March 2024



# Table of contents

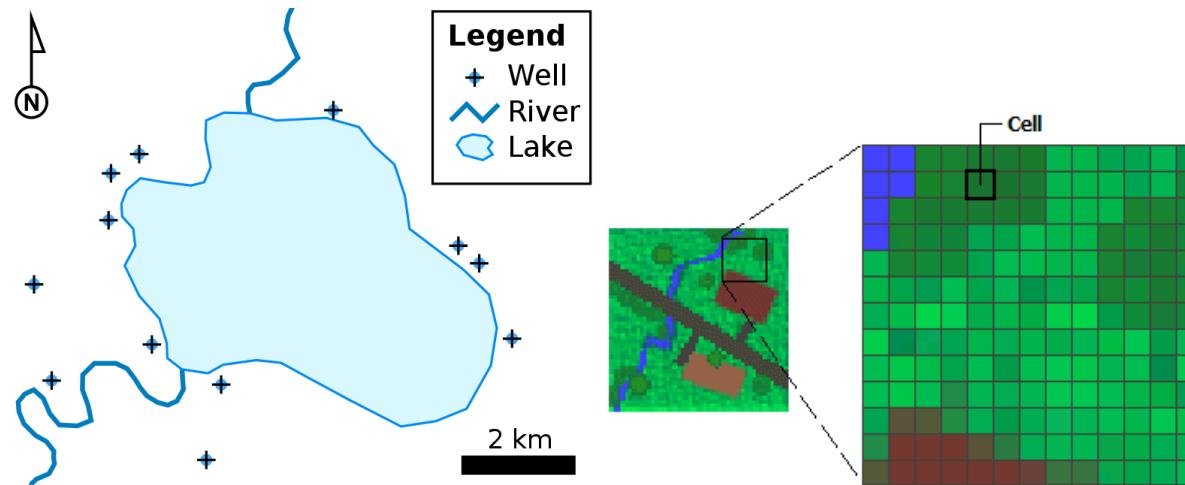
1. Overview of GIS concepts
2. Load and explore polygons, polylines, and points
3. Static maps
4. Interactive maps
5. Spatial operations applied on one dataset
6. Spatial operations applied on multiple datasets

# Overview of GIS concepts

**Spatial data:** The two main types of spatial data are **vector data** and **raster data**

## Vector data

- Points, lines, or polygons
- Common file formats include shapefiles (.shp) and geojsons (.geojson)
- Examples: polygons on countries, polylines of roads, points of schools

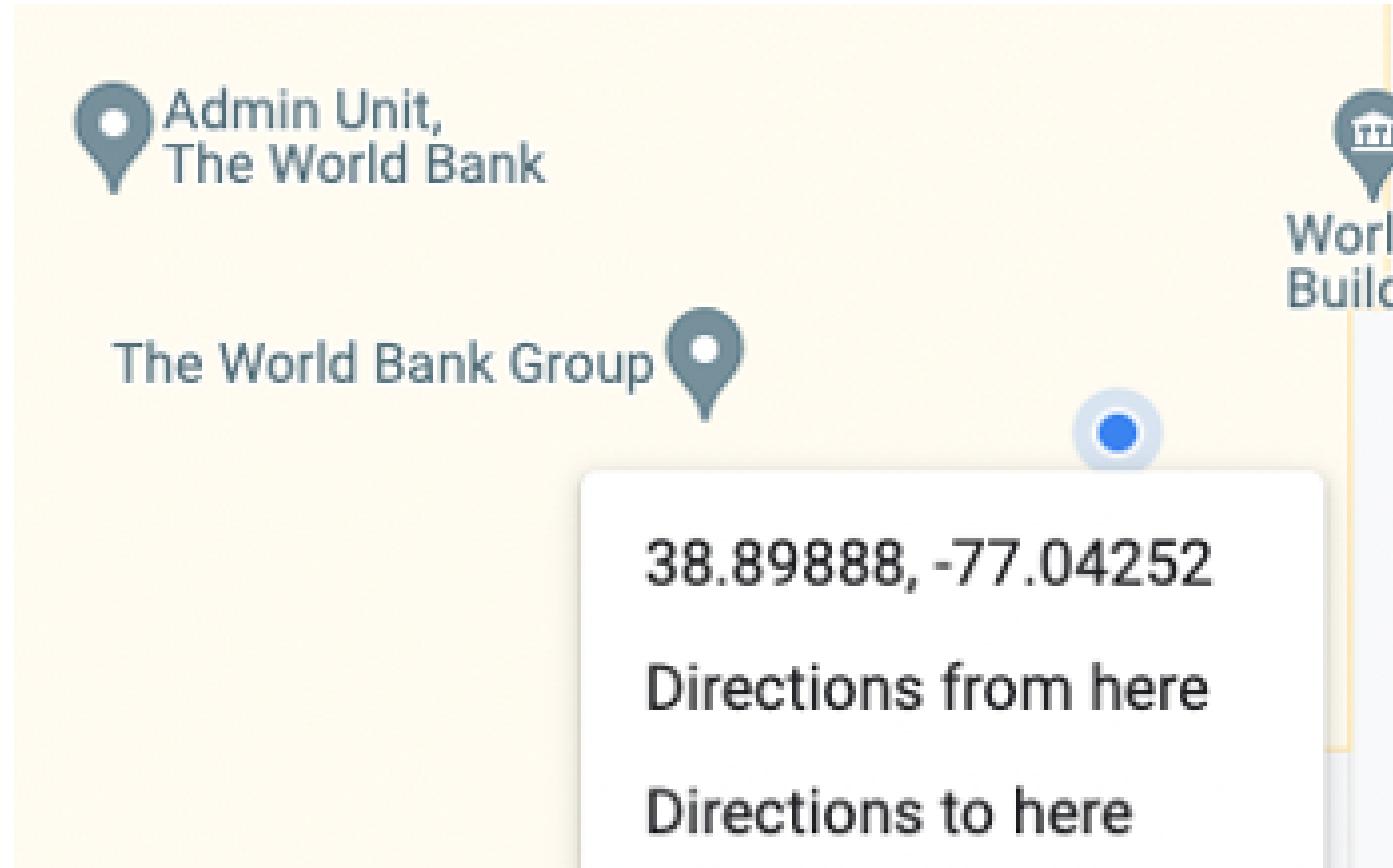


## Raster data

- Spatially referenced grid
- Common file format is a geotif (.tif)
- Example: Satellite imagery of nighttime lights

# Coordinate Reference Systems (CRS)

- Coordinate reference systems use pairs of numbers to define a location on the earth
- For example, the World Bank is at a latitude of 38.89 and a longitude of -77.04



# Coordinate Reference Systems (CRS)

There are many different coordinate reference systems, which can be grouped into **geographic** and **projected** coordinate reference systems. Geographic systems live on a sphere, while projected systems are “projected” onto a flat surface.

**Geographic (Sphere)**



40°W, 40°N

**Projected (Flat)**

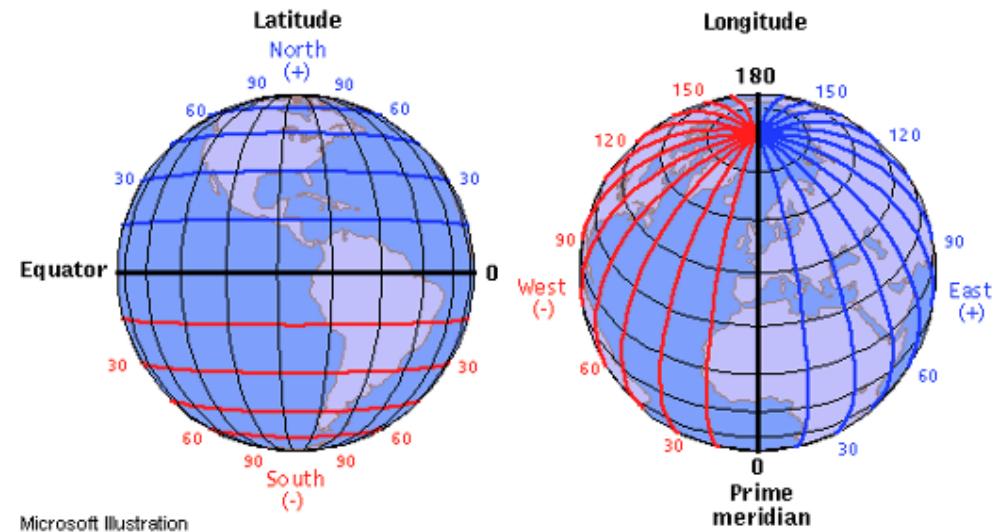


# Geographic Coordinate Systems

**Units:** Defined by latitude and longitude, which measure angles and units are typically in decimal degrees. (Eg, angle is latitude from the equator).

## Latitude & Longitude:

- On a grid X = longitude, Y = latitude; sometimes represented as (longitude, latitude).
- Also has become convention to report them in alphabetical order: (latitude, longitude) — such as in Google Maps.
- Valid range of latitude: -90 to 90
- Valid range of longitude: -180 to 180
- **{Tip}** Latitude sounds (and looks!) like latter.

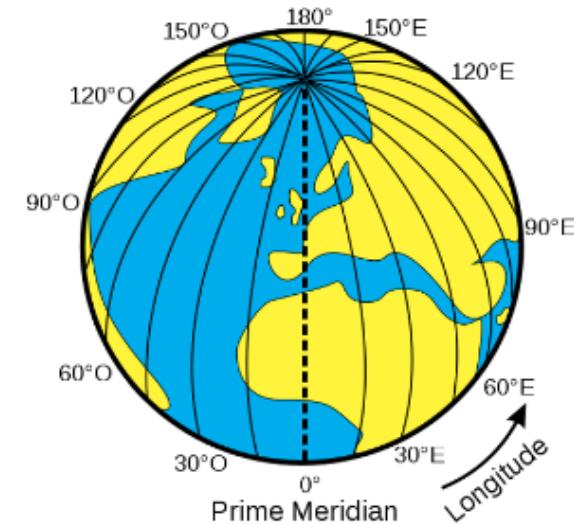


Microsoft Illustration

# Geographic Coordinate Systems

## Distance on a sphere

- At the equator (latitude = 0), a 1 decimal degree longitude distance is about 111km; towards the poles (latitude = -90 or 90), a 1 decimal degree longitude distance converges to 0 km.
- We must be careful (ie, use algorithms that account for a spherical earth) to calculate distances! The distance along a sphere is referred to as a **great circle distance**.
- Multiple options for spherical distance calculations, with trade-off between accuracy & complexity. (See distance section for details).



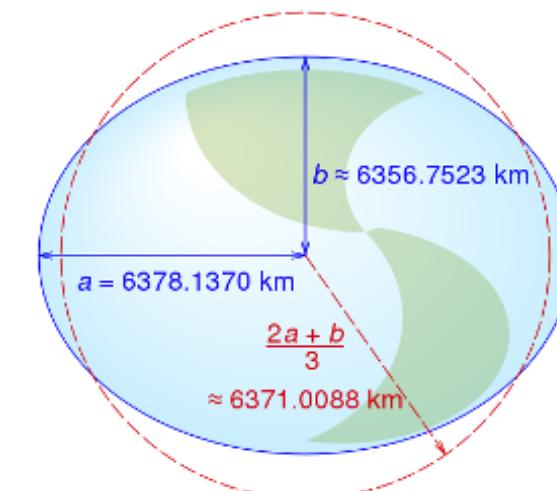
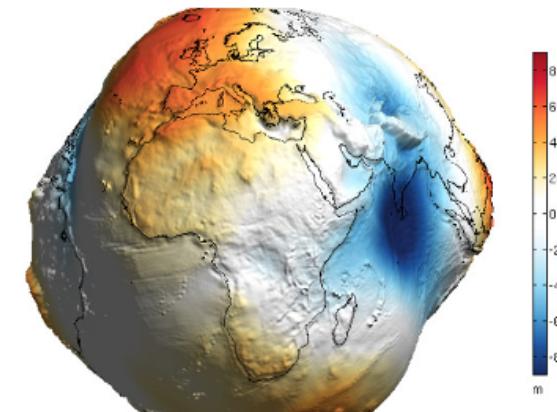
# Geographic Coordinate Systems

## Datums

- **Is the earth flat?** No!
- **Is the earth a sphere?** No!
- **Is the earth a lumpy ellipsoid?** Yes!

The earth is a lumpy ellipsoid, a bit flattened at the poles.

- A **datum** is a model of the earth that is used in mapping. One of the most common datums is **WGS 84**, which is used by the Global Positional System (GPS).
- A datum is a reference ellipsoid that approximates the shape of the earth.
- Other datums exist, and the latitude and longitude values for a specific location will be different depending on the datum.



# Projected Coordinate Systems

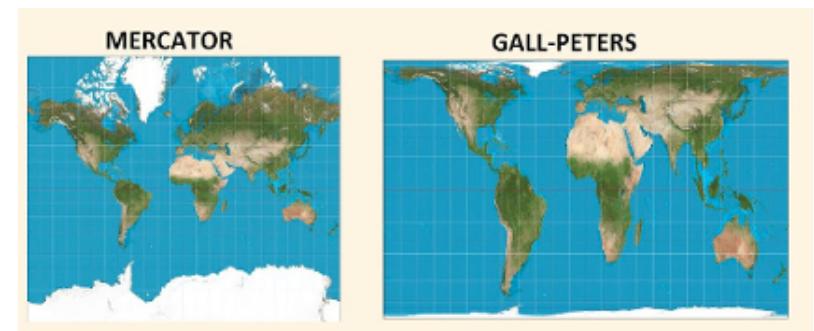
Projected coordinate systems project spatial data from a 3D to 2D surface.

**Distortions:** Projections will distort some combination of distance, area, shape or direction. Different projections can minimize distorting some aspect at the expense of others.

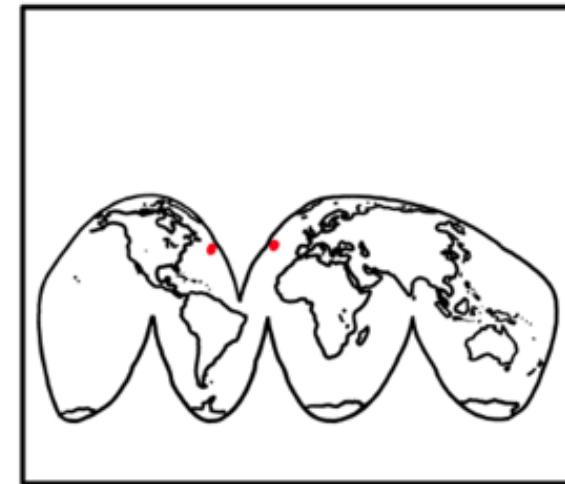
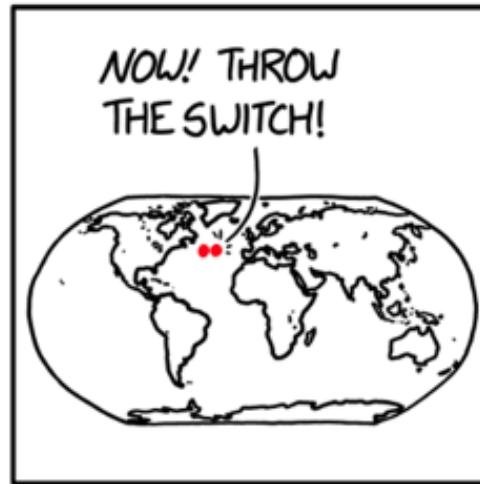
**Units:** When projected, points are represented as “northing” and “eastings.” Values are often represented in meters, where northings/eastings are the meter distance from some reference point. Consequently, values can be very large!

**Datums still relevant:** Projections start from some representation of the earth. Many projections (eg, [UTM](#)) use the WGS84 datum as a starting point (ie, reference datum), then project it onto a flat surface.

Click [here](#) to see why Toby & CJ are confused (hint: projections!)



# Projected Coordinate Systems



# Referencing coordinate reference systems

- There are many ways to reference coordinate systems, some of which are verbose.
- **PROJ** (Library for projections) way of referencing WGS84 `+proj=longlat  
+datum=WGS84 +no_defs +type=crs`
- **EPSG** Assigns numeric code to CRSs to make it easier to reference. Here, WGS84 is `4326`.

# Coordinate Reference Systems

Whenever have spatial data, need to know which coordinate reference system (CRS) the data is in.

- You wouldn't say "**I am 5 away**"
- You would say "**I am 5 [miles / kilometers / minutes / hours] away**" (units!)
- Similarly, a "complete" way to describe location would be: I am at **6.51 latitude, 3.52 longitude using the WGS 84 CRS**

# Introduction

- This session could be a whole course on its own, but we only have an hour and half.
- To narrow our subject, we will focus on only one type of spatial data, vector data.
- This is the most common type of spatial data that non-GIS experts will encounter in their work.
- We will use the `sf` package, which is the tidyverse-compatible package for geospatial data in R.
- For visualizing, we'll rely on `ggplot2` for static maps and `leaflet` for interactive maps

# Setup

1. Copy/paste the following code into a new RStudio script, **replacing "YOURFOLDERPATHHERE" with the folder within which you'll place this R project:**

```
library(usethis)
use_course(
  url = "https://github.com/worldbank/dime-r-training/archive/main.zip",
  destdir = "YOURFOLDERPATHHERE"
)
```

2. In the console, type in the requisite number to delete the .zip file (we don't need it anymore).
3. A new RStudio environment will open. Use this for the session today.

# Setup

Install new packages

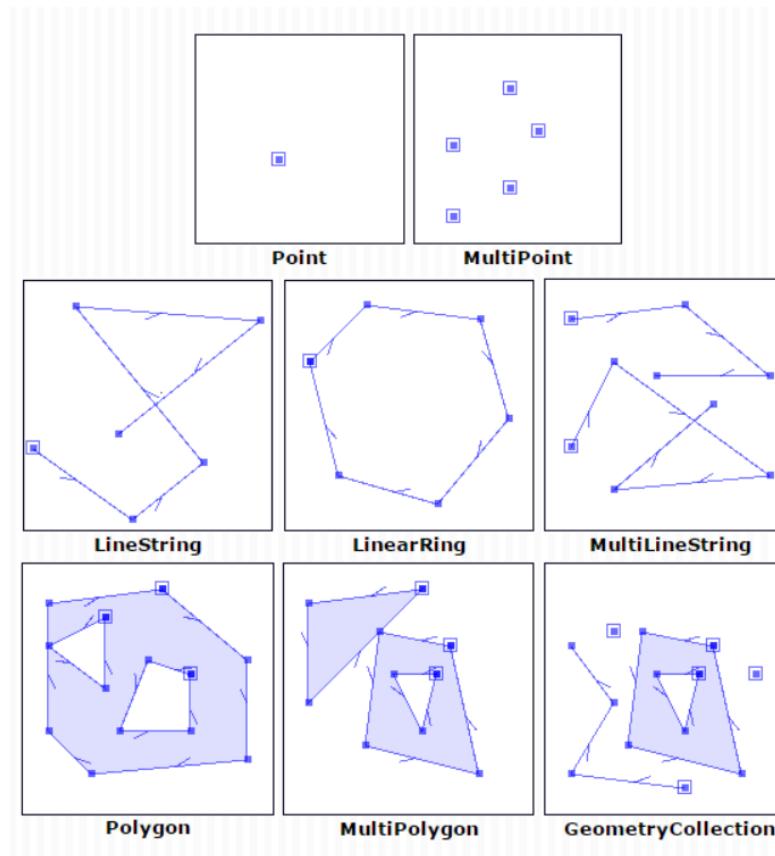
```
install.packages(c("sf",
                    "leaflet",
                    "geosphere"),
                    dependencies = TRUE)
```

And load them

```
library(here)
library(tidyverse)
library(sf)      # Simple features
library(leaflet) # Interactive map
library(geosphere) # Great circle distances
```

# Load and explore polylines, polylines, and points

The main package we'll rely on is the `sf` (simple features) package. With `sf`, spatial data is structured similarly to a **dataframe**; however, each row is associated with a **geometry**. Geometries can be one of the below types.



# Load and explore polygon

The first thing we will do in this session is to recreate this data set:

```
nga_sf <-  
  st_read(here("DataWork",  
              "DataSets",  
              "Final",  
              "gadm_nga_2.geojson"))  
  
## Reading layer `gadm_nga_2' from data source  
##   `C:/WBG/Repos/dime-r-training/DataWork/DataSets/Final/gadm_nga_2.geojson'  
##   using driver 'GeoJSON'  
## Simple feature collection with 775 features and 12 fields  
## Geometry type: MULTIPOLYGON  
## Dimension:      XY  
## Bounding box:  xmin: 2.668431 ymin: 4.270418 xmax: 14.67642 ymax: 13.89201  
## Geodetic CRS:  WGS 84
```

# Exploring the data

Look at first few observations

```
head(nga_sf)
```

```
## Simple feature collection with 6 features and 12 fields
## Geometry type: MULTIPOLYGON
## Dimension: XY
## Bounding box: xmin: 7.291035 ymin: 5.037226 xmax: 7.962045 ymax: 5.837777
## Geodetic CRS: WGS 84
##   ID_0 COUNTRY NAME_1 NL_NAME_1      ID_2          NAME_2 VARNAME_2 NL_NAME_2
## 1 NGA Nigeria Abia      NGA.1.1_1      Aba North
## 2 NGA Nigeria Abia      NGA.1.2_1      Aba South
## 3 NGA Nigeria Abia      NGA.1.3_1 Aiochukw Aiochukwu
## 4 NGA Nigeria Abia      NGA.1.4_1      Bende
## 5 NGA Nigeria Abia      NGA.1.5_1     Ikwuano
## 6 NGA Nigeria Abia      NGA.1.6_1 Isiala Ngwa North
##           TYPE_2    ENGTTYPE_2 CC_2 HASC_2                      geometry
## 1 Local Authority Local Authority MULTIPOLYGON (((7.378114 5.....
## 2 Local Authority Local Authority MULTIPOLYGON (((7.314784 5.....
## 3 Local Authority Local Authority MULTIPOLYGON (((7.641509 5.....
## 4 Local Authority Local Authority MULTIPOLYGON (((7.635005 5.....
```

# Exploring the data

Number of rows

```
nrow(nga_sf)
```

```
## [1] 775
```

# Exploring the data

Check coordinate reference system

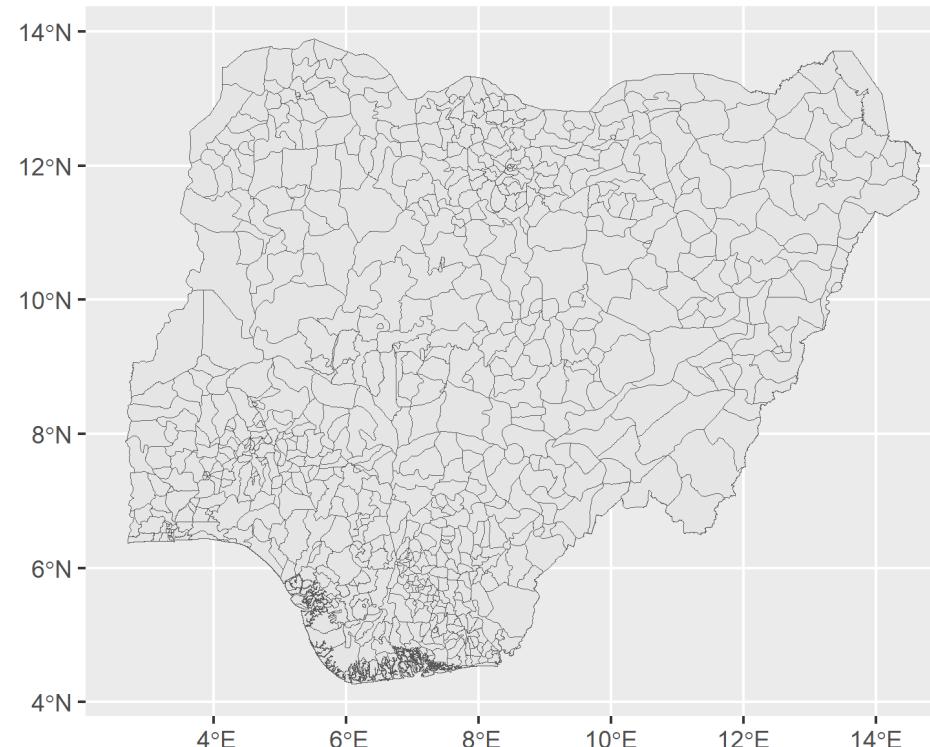
```
st_crs(nga_sf)
```

```
## Coordinate Reference System:  
##   User input: WGS 84  
##   wkt:  
## GEOCRS["WGS 84",  
##         DATUM["World Geodetic System 1984",  
##                 ELLIPSOID["WGS 84",6378137,298.257223563,  
##                             LENGTHUNIT["metre",1]],  
##                 PRIMEM["Greenwich",0,  
##                           ANGLEUNIT["degree",0.0174532925199433]],  
##                 CS[ellipsoidal,2],  
##                   AXIS["geodetic latitude (Lat)",north,  
##                         ORDER[1],  
##                         ANGLEUNIT["degree",0.0174532925199433]],  
##                   AXIS["geodetic longitude (Lon)",east,  
##                         ORDER[2],  
##                         ANGLEUNIT["degree",0.0174532925199433]],  
##                 ID["EPSG",4326]]
```

# Exploring the data

Plot the data. To plot using `ggplot2`, we use the `geom_sf` geometry.

```
ggplot() +  
  geom_sf(data = nga_sf)
```



# Attributes of data

We want the area of each location, but we don't have a variable for area

```
names(nga_sf)
```

```
## [1] "ID_0"      "COUNTRY"    "NAME_1"     "NL_NAME_1"  "ID_2"      "NAME_2"  
## [7] "VARNAME_2"  "NL_NAME_2"   "TYPE_2"     "ENGTTYPE_2" "CC_2"      "HASC_2"  
## [13] "geometry"
```

# Attributes of data

Determine area. Note the CRS is spherical (WGS84), but `st_area` gives area in meters squared. R uses s2 geometry for this.

```
st_area(nga_sf)
```

```
## Units: [m^2]
## [1] 18115918 51130486 386126742 603925481 237940893 253475003
## [7] 261346460 336386622 374567649 568416816 192717817 106283762
## [13] 246502187 293336257 451959706 231141018 131269420 2289574417
## [19] 4268043310 1884615977 1865638766 1084336031 699838559 2732234473
## [25] 1725588389 1167047612 585132151 508809998 2175752242 742913573
## [31] 654421163 209951938 813298273 1605421682 4324634661 4584293228
## [37] 139996668 370270051 177485439 96405062 189114803 139599636
## [43] 299011048 193538372 171492891 210962698 175747588 316253840
## [49] 130882266 269754902 334120006 133544979 473185832 255878306
## [55] 243461164 346695186 131733298 119971004 204247037 232632892
## [61] 275252661 151590665 81667384 470353728 53736209 279593007
## [67] 291550813 127666912 188377399 162798931 380294600 743228061
## [73] 105322234 389635592 171446836 570044307 58595577 107371646
## [79] 87655555 111134295 222332919 94111356 64805084 168774033
## [85] 398905868 28143090 17159027 364759033 231919200 131681287
## [91] 6133937818 3976854006 653949021 988289702 2498187950 528634061
```

# Operations similar to dataframes

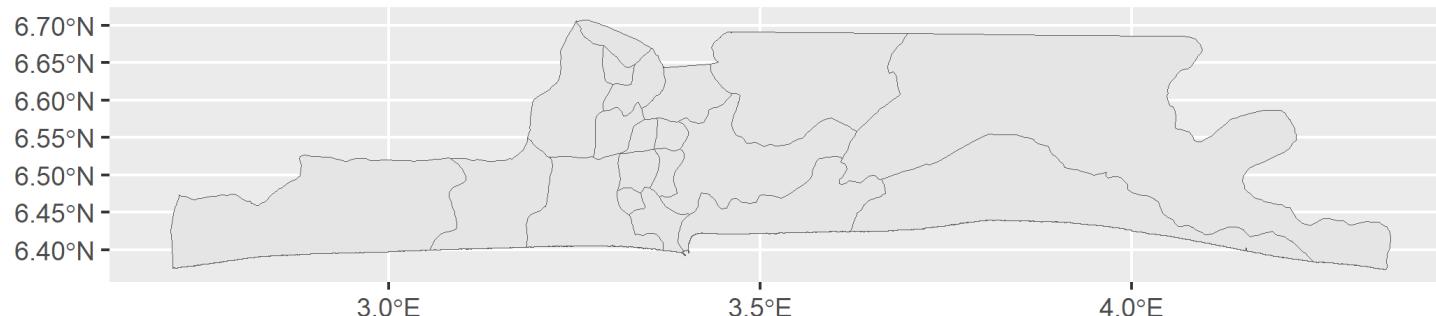
Create new dataset that is just lagos

```
lagos_sf <- nga_sf %>%  
  filter(NAME_1 == "Lagos")
```

# Operations similar to dataframes

Plot the dataframe

```
ggplot() +  
  geom_sf(data = lagos_sf)
```



# Load and explore polyline

## Exercise:

- Load the roads data `osm_lagos_roads.geojson` and name the object `roads_sf`
- Look at the first few observations
- Check the coordinate reference system
- Map the polyline

## Solution:

```
roads_sf <- st_read(here("DataWork", "DataSets", "Final", "osm_lagos_roads.geojson"))

head(roads_sf)

st_crs(roads_sf)

ggplot() +
  geom_sf(data = roads_sf)
```

02 : 00

# Load and explore polyline

```
roads_sf <-  
  st_read(here("DataWork",  
              "DataSets",  
              "Final",  
              "osm_lagos_roads.geojson"))  
  
## Reading layer `osm_lagos_roads' from data source  
##   `C:\WBG\Repos\dime-r-training\DataWork\DataSet\Final\osm_lagos_roads.geojson'  
##   using driver `GeoJSON'  
## Simple feature collection with 3007 features and 34 fields  
## Geometry type: LINESTRING  
## Dimension:      XY  
## Bounding box:  xmin: 2.705985 ymin: 6.379386 xmax: 4.217506 ymax: 6.832435  
## Geodetic CRS:  WGS 84  
  
ggplot() +  
  geom_sf(data = roads_sf)
```

# Load and explore polyline

**Exercise:** Determine length of each line (hint: use `st_length`)

**Solution:**

```
st_length(roads_sf)
```

```
## Units: [m]
##      [1] 794.082050 238.905024 33555.236573 3528.646684    78.305014
##      [6] 1906.248364 101.261589   625.342359 130.178118 335.860392
##     [11] 63.894351   29.269625 336.744791 64.698424 5675.070068
##     [16] 811.315023 60.027144 13110.313140 488.358816 2547.259226
##     [21] 66.730590 56.824559 1372.839939 14361.919350 16510.001084
##     [26] 4298.850486 991.619944 180.578417 66.128111 664.387693
##     [31] 623.150141 39.472571 117.079424 51.801118 46.843432
##     [36] 44.336573 11.152053 1750.478856 2135.870227 3854.862119
##     [41] 3858.886958 2681.089557 2689.057651 1915.872778 5485.677899
##      [46] 1506.045111 1499.337941 1884.624020 192.656690 5491.550596
##      [51] 3553.259716 3546.534051 71.313887 2138.421303 671.348526
##      [56] 1793.092780 32.819448 189.943926 193.794458 312.942975
```

01:00

# Load and explore point data

We'll load a dataset of the location of schools in Lagos

```
schools_df <-  
  read_csv(here("DataWork",  
               "DataSets",  
               "Final",  
               "osm_lagos_schools.csv"))  
  
## Rows: 429 Columns: 33  
## — Column specification ——————  
## Delimiter: ","  
## chr (28): element_type, amenity, min_age, name, opening_hours, wheelchair, ...  
## dbl (4): osmid, addr:postcode, latitude, longitude  
## date (1): survey:date  
##  
## i Use `spec()` to retrieve the full column specification for this data.  
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

# Explore data

```
head(schools_df)
```

```
## # A tibble: 6 × 33
##   element_type      osmid amenity min_age name  opening_hours wheelchair website
##   <chr>          <dbl> <chr>    <chr>  <chr>       <chr>     <chr>
## 1 node            8.53e 9 school  3yrs   St. ... Mo-Fr 07:45 ... yes    <NA>
## 2 way             3.79e 8 school  <NA>    Alta... <NA>       <NA>    https:...
## 3 way             5.59e 8 school  <NA>    Nean... <NA>       <NA>    <NA>
## 4 node            9.55e 9 school  <NA>    Bich... <NA>       <NA>    http:/...
## 5 node            1.11e10 school <NA>    Chri... <NA>       <NA>    <NA>
## 6 node            1.11e10 school <NA>    Cran... <NA>       <NA>    <NA>
## # i 25 more variables: wikidata <chr>, wikipedia <chr>, `addr:city` <chr>,
## # `addr:postcode` <dbl>, `addr:street` <chr>, email <chr>, phone <chr>,
## # building <chr>, grades <chr>, `addr:housenumber` <chr>,
## # official_name <chr>, short_name <chr>, `survey:date` <date>,
## # description <chr>, source <chr>, `name:en` <chr>, operator <chr>,
## # religion <chr>, denomination <chr>, `operator:type` <chr>,
## # `isced:level` <chr>, landuse <chr>, osm_type <chr>, latitude <dbl>, ...
```

# Explore data

```
names(schools_df)
```

```
## [1] "element_type"      "osmid"          "amenity"        "min_age"  
## [5] "name"              "opening_hours"   "wheelchair"     "website"  
## [9] "wikidata"           "wikipedia"      "addr:city"      "addr:postcode"  
## [13] "addr:street"        "email"          "phone"          "building"  
## [17] "grades"             "addr:housenumber" "official_name"   "short_name"  
## [21] "survey:date"         "description"    "source"         "name:en"  
## [25] "operator"            "religion"       "denomination"  "operator:type"  
## [29] "isced:level"         "landuse"        "osm_type"       "latitude"  
## [33] "longitude"
```

# Convert to spatial object

We define the (1) coordinates (longitude and latitude) and (2) CRS. **Note:** We must determine the CRS from the data metadata. This dataset comes from OpenStreetMaps, which uses WGS:4326.

**Assigning the incorrect CRS is one of the most common sources of issues I see with geospatial work. If something looks weird, check the CRS!**

```
schools_sf <- st_as_sf(schools_df,  
                        coords = c("longitude", "latitude"),  
                        crs = 4326)
```

# Convert to spatial object

```
head(schools_sf$geometry)

## Geometry set for 6 features
## Geometry type: POINT
## Dimension: XY
## Bounding box: xmin: 3.62856 ymin: 6.470489 xmax: 4.0054 ymax: 6.6652
## Geodetic CRS: WGS 84
## First 5 geometries:

## POINT (4.0054 6.6652)

## POINT (3.96229 6.616529)

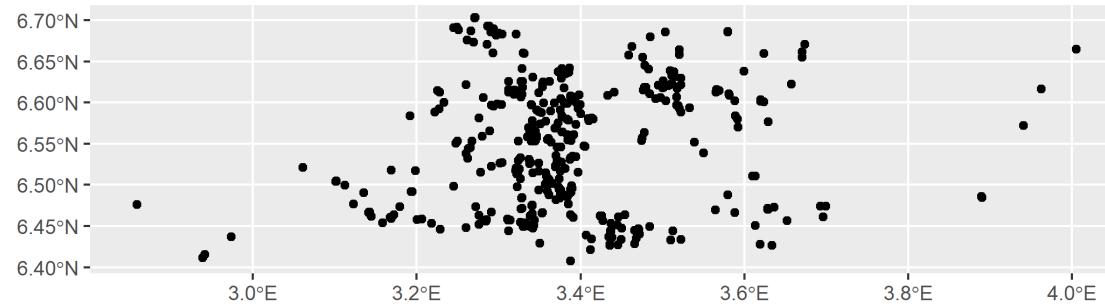
## POINT (3.940881 6.572413)

## POINT (3.692594 6.474112)

## POINT (3.629154 6.470489)
```

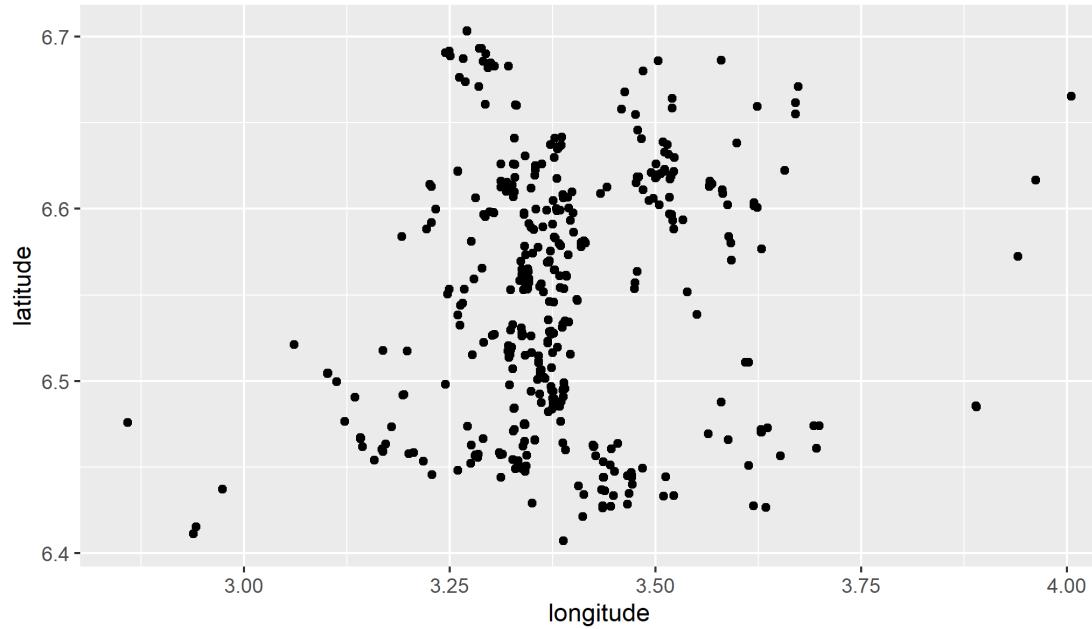
# Map points object: Using sf

```
ggplot() +  
  geom_sf(data = schools_sf)
```



# Map points object: Using dataframe

```
ggplot() +  
  geom_point(data = schools_df,  
             aes(x = longitude,  
                  y = latitude))
```



# Map points objects

**Question:** Why do the maps look different? Map using `sf` looks a bit squished!

**Solution:**

- Units are in decimal degrees.
- The length (eg, meters) between degrees of longitude shrinks as we go towards the equator.
- `geom_point` thinks 1 degree distance is the same for latitudes and longitudes (dumb!).
- `geom_sf` knows that we are in `WGS:4326` (spherical CRS), and adjusts map to minimize distortions (smart!).

00 : 30

# Make better static map

Lets make a better static map.

```
# Adding a variable with squared km
lagos_sf <- lagos_sf %>%
  mutate(area_m = lagos_sf %>% st_area() %>% as.numeric(),
        area_km = area_m / 1000^2)

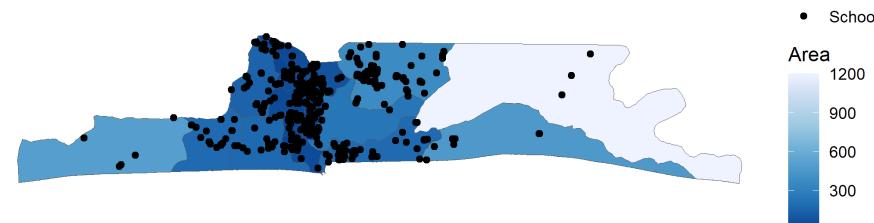
# Plotting
ggplot() +
  geom_sf(data = lagos_sf,
          aes(fill = area_km)) +
  labs(fill = "Area") +
  scale_fill_distiller(palette = "Blues") +
  theme_void()
```



# Make better static map

Lets add another spatial layer

```
ggplot() +  
  geom_sf(data = lagos_sf,  
           aes(fill = area_km)) +  
  geom_sf(data = schools_sf,  
           aes(color = "Schools")) +  
  labs(fill = "Area",  
       color = NULL) +  
  scale_fill_distiller(palette = "Blues") +  
  scale_color_manual(values = "black") +  
  theme_void()
```



# Interactive map

We use the `leaflet` package to make interactive maps. Leaflet is a JavaScript library, but the `leaflet` R package allows making interactive maps using R. Use of leaflet somewhat mimics how we use ggplot.

- Start with `leaflet()` (instead of `ggplot()`)
- Add spatial layers, defining type of layer (similar to geometries)

```
leaflet() %>%  
  addTiles() # Basemap
```

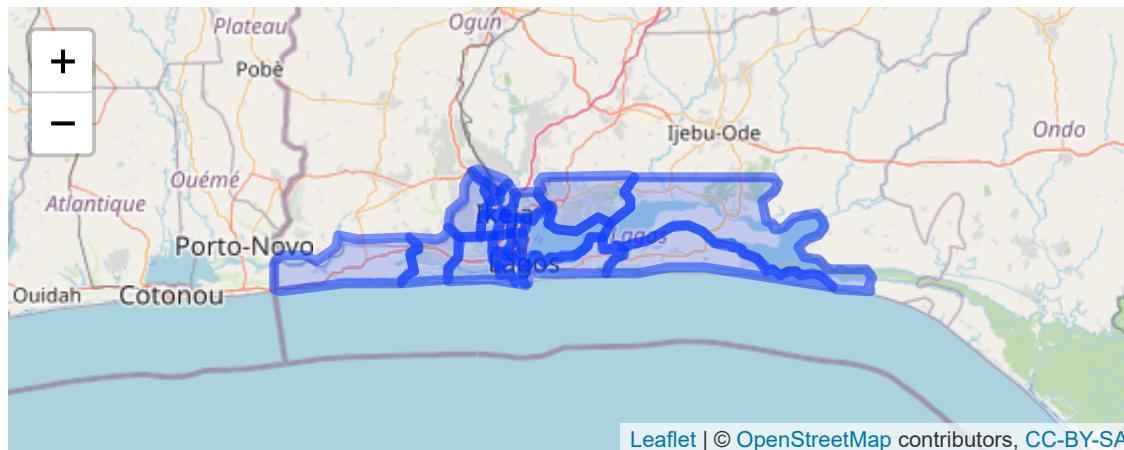


# Interactive map

We use the `leaflet` package to make interactive maps. Leaflet is a JavaScript library, but the `leaflet` R package allows making interactive maps using R. Use of leaflet somewhat mimics how we use ggplot.

- Start with `leaflet()` (instead of `ggplot()`)
- Add spatial layers, defining type of layer (similar to geometries)

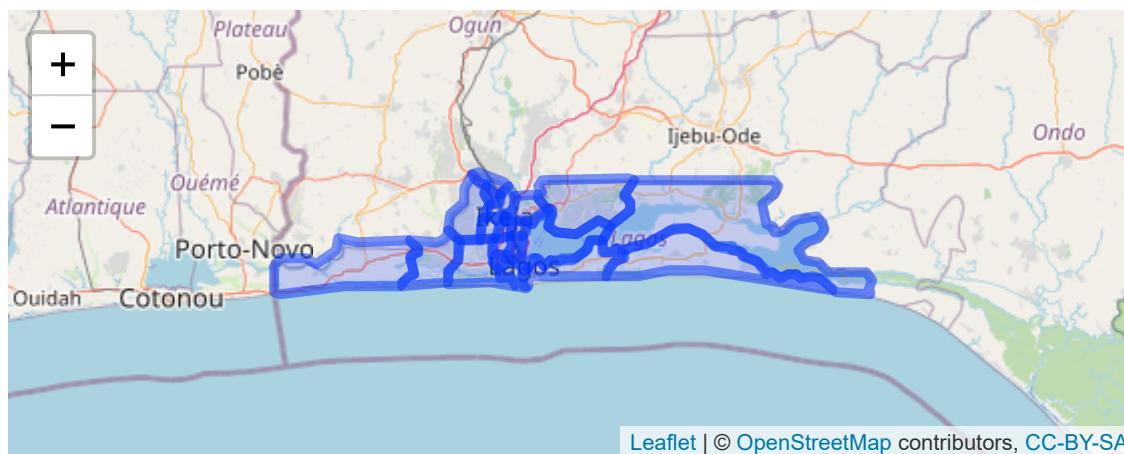
```
leaflet() %>%  
  addTiles() %>%  
  addPolygons(data = lagos_sf)
```



# Interactive map

Add a pop-up

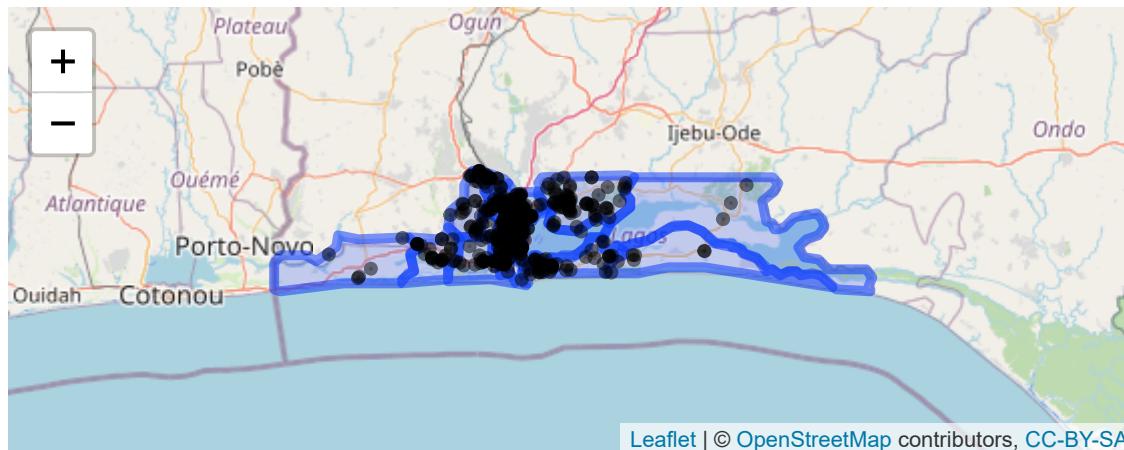
```
leaflet() %>%  
  addTiles() %>%  
  addPolygons(data = lagos_sf,  
              popup = ~NAME_2)
```



# Interactive map

Add more than one layer

```
leaflet() %>%  
  addTiles()  
  addPolygons(data = lagos_sf,  
              popup = ~NAME_2)  
  addCircles(data = schools_sf,  
             popup = ~name,  
             color = "black")
```



# Interactive map of roads

**Exercise:** Create a leaflet map with roads, using the `roads_sf` dataset. (**Hint:** Use `addPolylines()`)

**Solution:**

```
leaflet() %>%  
  addTiles() %>%  
  addPolylines(data = roads_sf)
```

# Interactive map of roads

```
leaflet() %>%  
  addTiles() %>%  
  addPolylines(data = roads_sf)
```



# Interactive maps

We can spent lots of time going over what we can do with leaflet - but that would take up too much time. [This resource](#) provides helpful tutorials for things like:

- Changing the basemap
- Adding colors
- Adding a legend
- And much more!

# Spatial operations applied on single dataset

- `st_transform`: Transform CRS
- `st_buffer`: Buffer point/line/polygon
- `st_combine`: Dissolve by attribute
- `st_convex_hull`: Create convex hull
- `st_centroid`: Create new sf object that uses the centroid
- `st_drop_geometry`: Drop geometry; convert from sf to dataframe
- `st_coordinates`: Get matrix of coordinates
- `st_bbox`: Get bounding box

# Transform CRS

We want to compute the length of roads in Lagos using the roads dataset. The roads dataset is currently in a geographic CRS (WGS84), where the units are in decimal degrees. We'll transform the CRS to a projected CRS that is suitable for Nigeria ([EPSG:32632](#)), and where the units will be in meters.

**Note that coordinate values are large!** Values are large because units are in meters. Large coordinate values suggest projected CRS; latitude is between -90 and 90 and longitude is between -180 and 180.

```
schools_utm_sf <- st_transform(schools_sf, 32632)

schools_utm_sf$geometry %>% head(2) %>% print()

## Geometry set for 2 features
## Geometry type: POINT
## Dimension:      XY
## Bounding box:  xmin: -57566.51 ymin: 734187.7 xmax: -52728.69 ymax: 739539.7
## Projected CRS: WGS 84 / UTM zone 32N

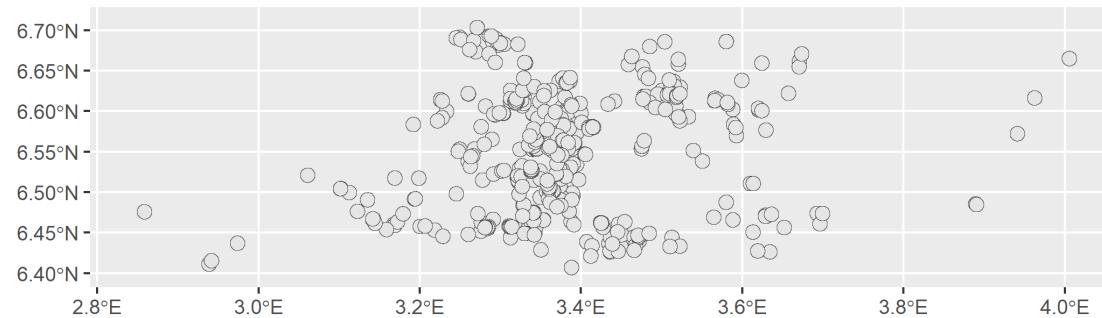
## POINT (-52728.69 739539.7)

## POINT (-57566.51 734187.7)
```

# Buffer

We have the points of schools. Now we create a 1km buffer around schools.

```
schools_1km_sf <- schools_sf %>%  
  st_buffer(dist = 1000) # Units are in meters. Thanks s2!  
  
ggplot() +  
  geom_sf(data = schools_1km_sf)
```

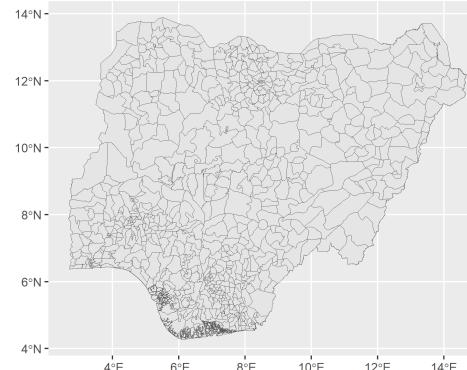


# Dissolve by an attribute

Below we have the second administrative regions of Nigeria. Using this dataset, let's create a new object at the first administrative region level.

```
nga1_sf <- nga_sf %>%
  group_by(NAME_1) %>%
  summarise(geometry = st_combine(geometry)) %>%
  ungroup()

ggplot() +
  geom_sf(data = nga1_sf)
```

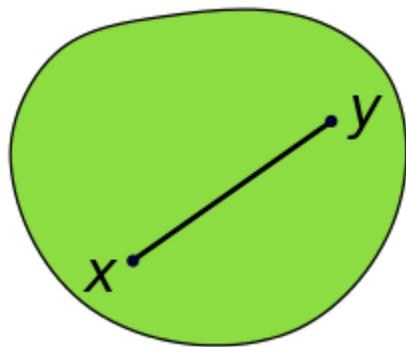


# Convex Hull

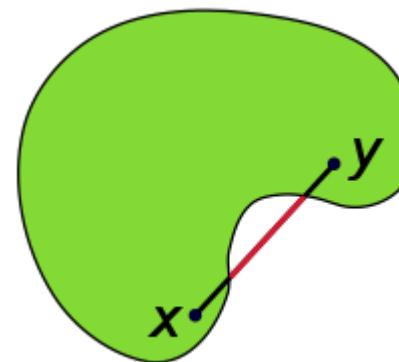
**Simple definition:** Get the outer-most coordinates of a shape and connect-the-dots.

**Formal definition:** A convex hull of a shape the smallest "convex set" that contains it. (A **convex set** is where a straight line can be drawn anywhere in the space and the space fully contains the line).

**Convex**



**Not convex**



**Source:** Wikipedia

# Convex hull

In the below example, we create a convex hull around schools; creating a polygon that includes all schools.

## Incorrect attempt

```
schools_chull_sf <- schools_sf %>%  
  st_convex_hull()  
  
nrow(schools_chull_sf)
```

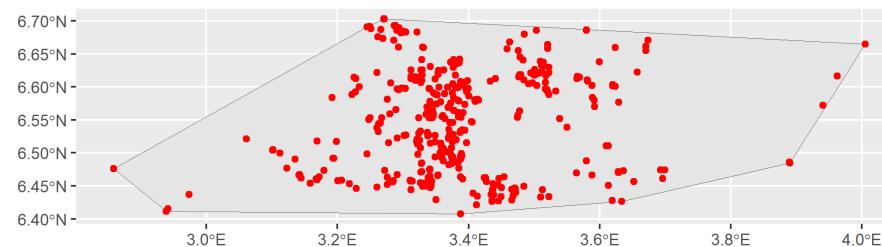
```
## [1] 429
```

# Convex hull

Correct

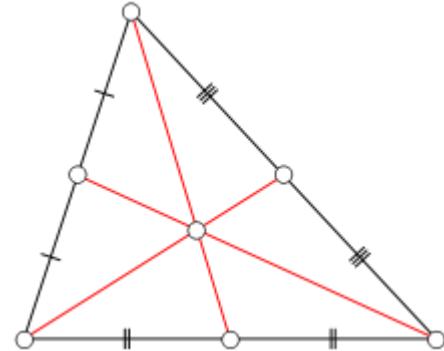
```
schools_chull2_sf <- schools_sf %>%
  summarise(geometry = st_combine(geometry)) %>%
  st_convex_hull()

ggplot() +
  geom_sf(data = schools_chull2_sf) +
  geom_sf(data = schools_sf, color = "red")
```



# Determine centroid

Sometimes we want to represent a polygon or polyline as a single point. For this, we can compute the centroid (ie, geographic center) of a polygon/polyline.



**Source:** Wikipedia

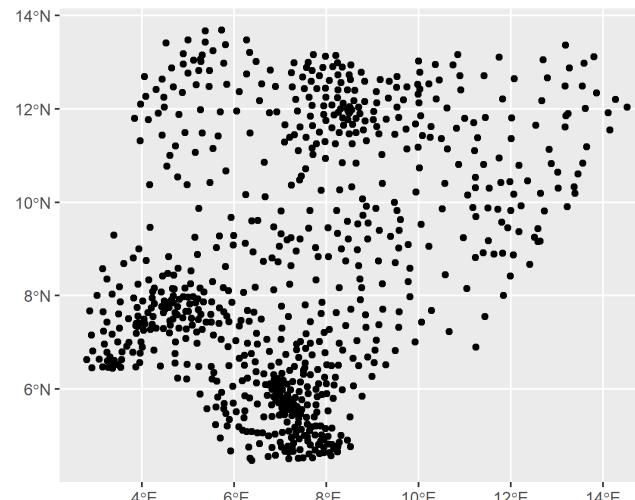
# Determine centroid

Determine centroid of Nigeria ADM2s

```
nga_c_sf <- st_centroid(nga_sf)
```

```
## Warning: st_centroid assumes attributes are constant over geometries
```

```
ggplot() +  
  geom_sf(data = nga_c_sf)
```



# Remove geometry

## Incorrect approach

```
lagos_sf %>%
  select(-geometry) %>%
  head()
```

```
## Simple feature collection with 6 features and 14 fields
## Geometry type: MULTIPOLYGON
## Dimension: XY
## Bounding box: xmin: 2.70628 ymin: 6.375416 xmax: 3.404933 ymax: 6.705109
## Geodetic CRS: WGS 84
##   ID_0 COUNTRY NAME_1 NL_NAME_1      ID_2          NAME_2 VARNAME_2 NL_NAME_2
## 1 NGA Nigeria Lagos      NGA.25.1_1      Agege
## 2 NGA Nigeria Lagos      NGA.25.2_1 Ajeromi/Ifelodun
## 3 NGA Nigeria Lagos      NGA.25.3_1     Alimosho
## 4 NGA Nigeria Lagos      NGA.25.4_1    Amuwo Odofin
## 5 NGA Nigeria Lagos      NGA.25.5_1      Apapa
## 6 NGA Nigeria Lagos      NGA.25.6_1    Badagary  Badagry
##           TYPE_2      ENGTTYPE_2 CC_2 HASC_2 area_m area_km
## 1 Local Authority Local Authority      16575387 16.57539
## 2 Local Authority Local Authority      11643619 11.64362
```

# Remove geometry

## Correct

```
lagos_sf %>%  
  st_drop_geometry() %>%  
  head()
```

```
##   ID_0 COUNTRY NAME_1 NL_NAME_1      ID_2          NAME_2 VARNAME_2 NL_NAME_2  
## 1 NGA Nigeria Lagos      NGA.25.1_1      Agege  
## 2 NGA Nigeria Lagos      NGA.25.2_1 Ajeromi/Ifeodun  
## 3 NGA Nigeria Lagos      NGA.25.3_1     Alimosho  
## 4 NGA Nigeria Lagos      NGA.25.4_1    Amuwo Odofin  
## 5 NGA Nigeria Lagos      NGA.25.5_1      Apapa  
## 6 NGA Nigeria Lagos      NGA.25.6_1    Badagary  Badagry  
##           TYPE_2      ENGTTYPE_2 CC_2 HASC_2      area_m      area_km  
## 1 Local Authority Local Authority      16575387  16.57539  
## 2 Local Authority Local Authority      11643619  11.64362  
## 3 Local Authority Local Authority      148062074 148.06207  
## 4 Local Authority Local Authority      174091716 174.09172  
## 5 Local Authority Local Authority      41688102  41.68810  
## 6 Local Authority Local Authority      494706196 494.70620
```

# Grab coordinates

Create a matrix of coordinates

```
schools_sf %>%
  st_coordinates() %>%
  head()
```

```
##           X         Y
## [1,] 4.005400 6.665200
## [2,] 3.962290 6.616529
## [3,] 3.940881 6.572413
## [4,] 3.692594 6.474112
## [5,] 3.629154 6.470489
## [6,] 3.628560 6.471747
```

# Get bounding box

```
schools_sf %>%
  st_bbox()
```

```
##      xmin      ymin      xmax      ymax
## 2.858534 6.407246 4.005400 6.703460
```

# Exercise

**Exercise:** Create a polyline of all trunk roads in Lagos, and buffer the polyline by 10 meters. In `roads_sf`, the `highway` variable notes road types.

## Solution:

```
roads_sf %>%
  filter(highway == "trunk") %>%
  summarise(geometry = st_combine(geometry)) %>%
  st_buffer(dist = 10)

## Simple feature collection with 1 feature and 0 fields
## Geometry type: MULTIPOLYGON
## Dimension: XY
## Bounding box: xmin: 2.705843 ymin: 6.37871 xmax: 3.520564 ymax: 6.716115
## Geodetic CRS: WGS 84
## #> #> geometry
#> #>   POLYGON (((3.255848 6....
```

02 : 00

# Spatial operations using multiple datasets

- `st_distance`: Calculate distances.
- `st_intersects`: Indicates whether simple features intersect.
- `st_intersection`: Cut one spatial object based on another.
- `st_difference`: Remove part of spatial object based on another.
- `st_join`: Spatial join (ie, add attributes of one dataframe to another based on location).

# Distances

For this example, we'll compute the distance between each school to a motorway.

```
motor_sf <- roads_sf %>%
  filter(highway == "motorway")

# Matrix: distance of each school to each motorway
dist_mat <- st_distance(schools_sf, motor_sf)

# Take minimum distance for each school
dist_mat %>% apply(1, min) %>% head()

## [1] 67892.77 62462.18 59452.39 31450.97 24430.84 24376.20
```

# Distances

There are multiple ways to calculate distances!

- **Great circle:** sf, by default, uses s2 to compute distance (in meters) when data has a geographic CRS
- **Great circle:** Other formulas beyond s2, such as Haversine, Vincenty, and Karney's method. See the [geosphere](#) and [geodist](#) packages. Vincenty is more precise than Haversine, and Karney's method is more precise than Vincenty's method. Greater precision comes with heavy computation. For more information, see [here](#).
- **Projected:** We can use a projected CRS, where units are in meters already.

# Distances

```
# s2  
st_distance(schools_sf[1,], schools_sf[2,]) %>%  
  as.numeric()  
  
## [1] 7208.358
```

```
# Nigeria-specific CRS  
schools_utm_sf <- st_transform(schools_sf, 32632)  
st_distance(schools_utm_sf[1,], schools_utm_sf[2,]) %>%  
  as.numeric()  
  
## [1] 7214.479
```

```
# World mercator  
schools_merc_sf <- st_transform(schools_sf, 3395)  
st_distance(schools_merc_sf[1,], schools_merc_sf[2,])  
  as.numeric()  
  
## [1] 7238.156
```

```
# Haversine  
distHaversine(  
  p1 = schools_sf[1,] %>% st_coordinates,  
  p2 = schools_sf[2,] %>% st_coordinates)
```

```
## [1] 7216.422
```

```
# Vincenty's method  
distVincentySphere(  
  p1 = schools_sf[1,] %>% st_coordinates,  
  p2 = schools_sf[2,] %>% st_coordinates)
```

```
## [1] 7216.422
```

```
# Karney's method  
distGeo(p1 = schools_sf[1,] %>% st_coordinates,  
        p2 = schools_sf[2,] %>% st_coordinates)
```

```
## [1] 7189.913
```

# Intersects

For this example we'll determine which of Lagos's second administrative divisions intersects with a motorway.

```
# Sparse matrix
st_intersects(lagos_sf, motor_sf) %>% print()

## Sparse geometry binary predicate list of length 20, where the predicate
## was `intersects'
## first 10 elements:
## 1: (empty)
## 2: 51, 54, 55, 56, 57, 60
## 3: (empty)
## 4: 57, 60
## 5: 21, 22, 28, 29, 32, 35, 36, 57, 58, 59, ...
## 6: (empty)
## 7: (empty)
## 8: 9, 19, 22, 69, 72
## 9: (empty)
## 10: (empty)
```

# Intersects

Take `max` (`FALSE` corresponds to 0 and `TRUE` corresponds to 1). So taking max will yeild if unit intersects with *any* motorway

```
# Matrix  
st_intersects(lagos_sf, motor_sf, sparse = F) %>%  
  apply(1, max) %>%  
  head()  
  
## [1] 0 1 0 1 1 0
```

# Intersection

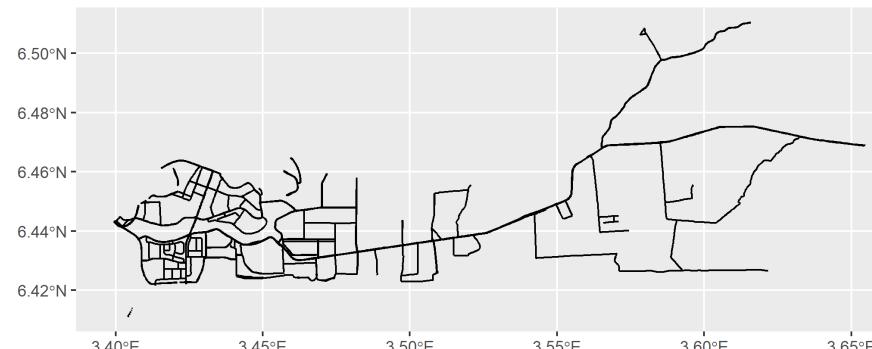
We have roads for all of Lagos. Here, we want to create new roads object that **only includes** roads in the Eti-Osa unit.

```
etis_sf <- lagos_sf %>%
  filter(NAME_2 == "Eti-Osa")

roads_etis_sf <- st_intersection(roads_sf, etis_sf)
```

```
## Warning: attribute variables are assumed to be spatially constant throughout
## all geometries
```

```
ggplot() +
  geom_sf(data = roads_etis_sf)
```



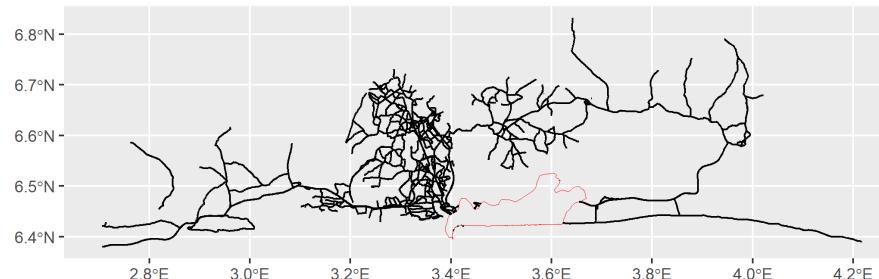
# Difference

We have roads for all of Lagos. Here, we want to create new roads object that **excludes** roads in the Eti-Osa unit.

```
roads_notetis_sf <- st_difference(roads_sf, etis_sf)

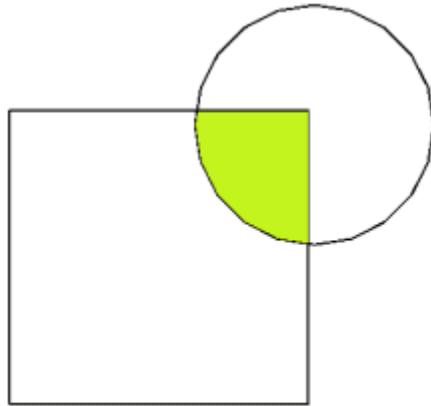
## Warning: attribute variables are assumed to be spatially constant throughout
## all geometries

ggplot() +
  geom_sf(data = etis_sf, fill = NA, color = "red") +
  geom_sf(data = roads_notetis_sf)
```

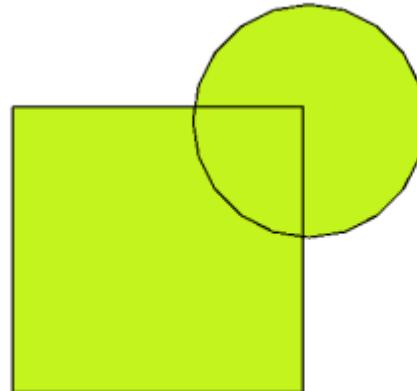


# Overlay

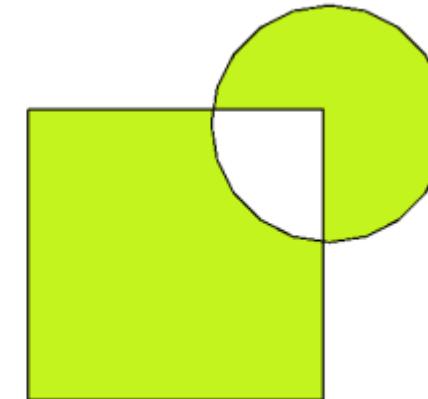
Intersections and differencing are **overlay** functions



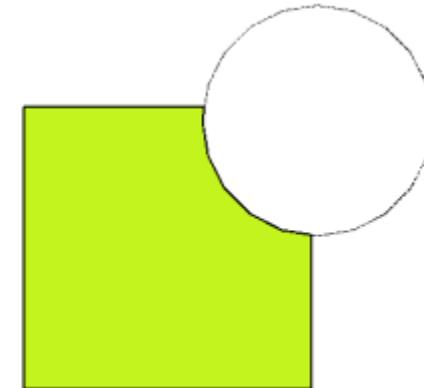
Intersection



Union



Symmetrical Difference



Difference

# Spatial join

We have a dataset of schools. The school dataframe contains information such as the school name, but not on the administrative region it's in. To add data on the administrative region that the school is in, we'll perform a spatial join.

Check the variable names. No names of second administrative division :(

```
names(schools_sf)
```

```
## [1] "element_type"      "osmid"           "amenity"          "min_age"
## [5] "name"              "opening_hours"    "wheelchair"       "website"
## [9] "wikidata"          "wikipedia"       "addr:city"        "addr:postcode"
## [13] "addr:street"       "email"            "phone"            "building"
## [17] "grades"            "addr:housenumber" "official_name"    "short_name"
## [21] "survey:date"       "description"     "source"           "name:en"
## [25] "operator"          "religion"         "denomination"   "operator:type"
## [29] "isced:level"       "landuse"          "osm_type"         "geometry"
```

# Spatial join

Use `st_join` to add attributes from `lagos_sf` to `schools_sf`. `st_join` is similar to other join methods (eg, `left_join`); instead of joining on a variable, we join based on location.

```
schools_lagos_sf <- st_join(schools_sf, lagos_sf)

schools_lagos_sf %>%
  names() %>%
  print() %>%
  tail(10)

## [1] "element_type"      "osmid"           "amenity"          "min_age"
## [5] "name"              "opening_hours"    "wheelchair"       "website"
## [9] "wikidata"          "wikipedia"       "addr:city"        "addr:postcode"
## [13] "addr:street"       "email"            "phone"            "building"
## [17] "grades"            "addr:housenumber" "official_name"    "short_name"
## [21] "survey:date"       "description"     "source"           "name:en"
## [25] "operator"          "religion"         "denomination"   "operator:type"
## [29] "isced:level"       "landuse"          "osm_type"         "geometry"
## [33] "ID_0"               "COUNTRY"          "NAME_1"           "NL_NAME_1"
## [37] "ID_2"               "NAME_2"            "VARNAME_2"        "NL_NAME_2"
## [41] "TYPE_2"             "ENGTYPE_2"        "CC_2"             "HASC_2"
## [45] "NAME_3"             "VARNAME_3"        "NL_NAME_3"        "HASC_3"
```

# Spatial join

**Exercise:** Make a static map using of Lagos administrative areas, where each administrative area polygon displays the number of schools.

**Solution:**

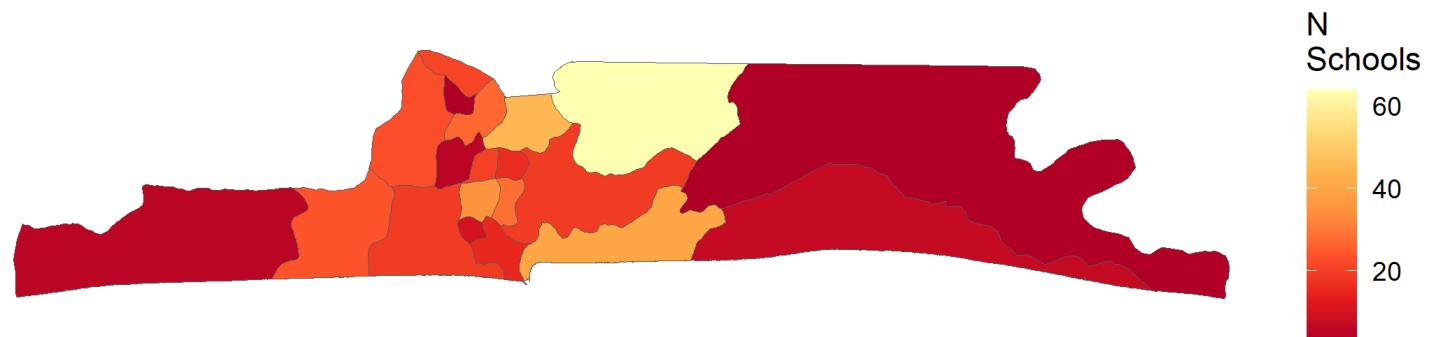
```
## Dataframe of number of schools per NAME_2
n_school_df <- schools_lagos_sf %>%
  st_drop_geometry() %>%
  group_by(NAME_2) %>%
  summarise(n_school = n()) %>%
  ungroup()

## Merge info with lagos_sf
lagos_sch_sf <- lagos_sf %>% left_join(n_school_df, by = "NAME_2")

## Map
p <- ggplot() +
  geom_sf(data = lagos_sch_sf,
          aes(fill = n_school))
```

# Spatial join

```
ggplot() +  
  geom_sf(data = lagos_sch_sf,  
          aes(fill = n_school)) +  
  labs(fill = "N\Schools") +  
  scale_fill_distiller(palette = "YlOrRd") +  
  theme_void()
```



# Resources

- sf package cheatsheet
- Spatial Data Science with Applications in R
- Geocomputation with R

# Thank you!

